FALCONAGENCY Oduriana

Introduction to Deep Learning

Poo Kuan Hoong

19th July 2016

Data Science Institute

- The Data Science Institute is a research center based in the Faculty of Computing & Informatics, Multimedia University.
- The members comprise of expertise across faculties such as Faculty of Computing and Informatics, Faculty of Engineering, Faculty of Management & Faculty of Information Science and Technology.
- Conduct research in leading data science areas including stream mining, video analytics, machine learning, deep learning, next generation data visualization and advanced data modelling.

Google DeepMind playing Atari Breakout

https://www.youtube.com/watch?v=V1eYniJ0Rnk

Google DeepMind Alphago

AlphaGo marks stark difference between AI and human intelligence

Daniel Susskind

Share · Author alerts · Print · Clio · Comments

Google's robot relied on processing power and data storage, writes Daniel Susskind

There are many ways of being smart that aren't smart like us." These are the words of Patrick Winston, a leading voice in the field of artificial intelligence. Although his idea is simple, its significance has been lost on most people thinking about the future of work. Yet this is the feature of AI that ought to preoccupy us the most.

From the 1950s to the 1980s, during the "first wave" of AI research, it was generally

What AlphaGo's Win Means for Your Job

COMMENTARY by Howard Yu @HowardHYu MARCH 21, 2016, 3:11 PM EDT

And the future of artificial intelligence

Just last week, machines crossed a momentous milestone. Google's AlphaGo, a computer algorithm, beat Go world champion Lee Sedol 4 to 1 in the ancient Chinese board game.

South Korean professional Go player Lee Se-Dol reviews the match after the fourth match against Google's artificial intelligence program, AlphaGo, during the Google DeepMind Challenge Match on March 13, 2016 in Seoul, South Korea.
Photograph by Google via Getty Images

RECOMMENDED FOR YOU

Unlike Western chess, which consists of about 40 turns in a

AlphaGo's Success Shows the Human Advantage Is Eroding Fast

Howard Yu is a professor of strategic management and innovation at IMD business school.

MARCH 9, 2016

It's impressive that computers can be programmed to conduct conversations that make them indistinguishable from a person. But Google's AlphaGo is demonstrating for the first time that machines can truly learn and think in a human way.

In 1996, IBM's Deep Blue program overwhelmed the world's greatest player, Garry Kasparov, with a brute force approach, in which the machine could account for all the possible outcomes.

AlphaGo's approach, on the other hand, is, potentially, a game changer. It can master games by adjusting as it goes.

Unlike chess, where each move affords about 40 options, the ancient board game Go, has up to 200 choices. The permutation of outcomes quickly compounds to a bewildering range of choices — more than the total number of atoms in the entire observable universe.

Baidu eye

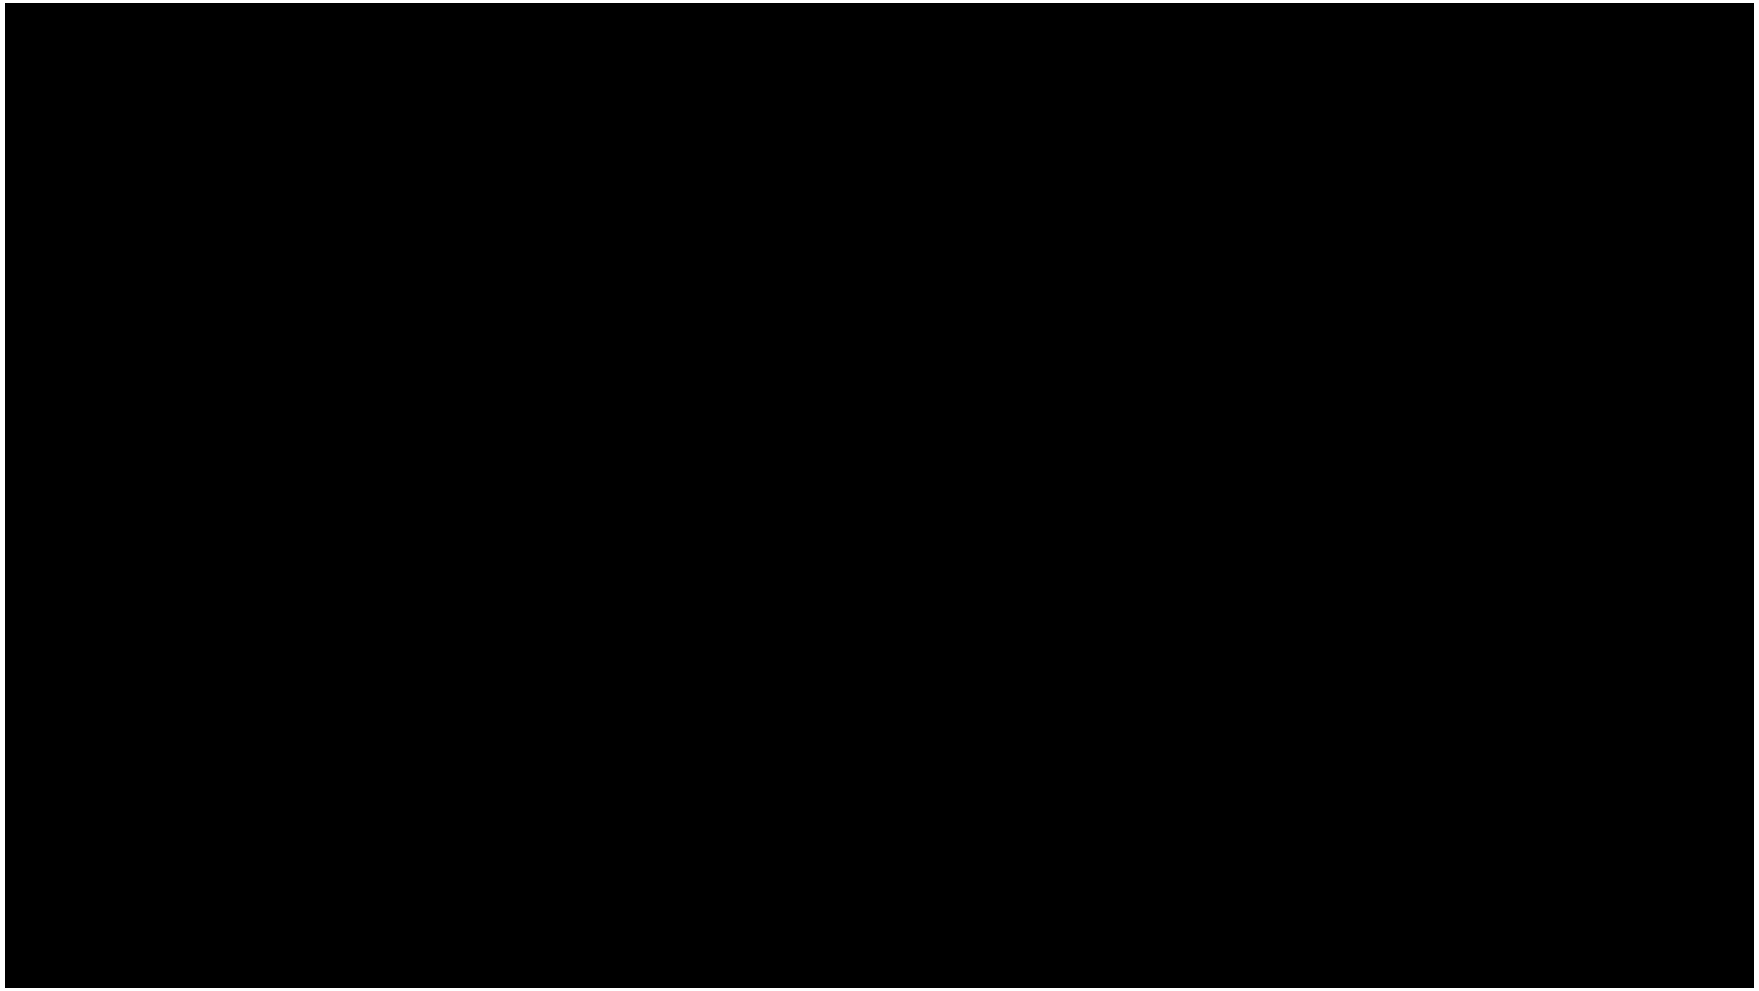

https://www.youtube.com/watch?v=Xe5RcJ1JY3c

Acknowledgement

Andrew Ng: Deep Learning, Self-Taught Learning and Unsupervised Feature Learning [Youtube]

Yann LeCun: Deep Learning Tutorial, ICML, Atlanta, 2013 [PDF]

Geoff Hinton, Yoshua Bengio & Yann LeCun: Deep Learning: NIPS2015 Tutorial [PDF]

Yoshua Bengio: Theano: A Python framework for fast computation of mathematical expressions. [URL]

Andrej Karpathy: Visualizing and Understanding Recurrent Networks, ICLR 2016, [PDF]

Outline

- A brief history of machine learning
- Understanding the human brain
- Neural Network: Concept, implementation and challenges
- Deep Belief Network (DBN): Concept and Application
- Convolutional Neural Network (CNN): Concept and Application
- Recurrent Neural Network (RNN): Concept and Application
- Deep Learning: Strengths, weaknesses and applications
- Deep Learning: Platforms, frameworks and libraries
- Demo

Introduction

- In the past 10 years, machine learning and Artificial Intelligence have shown tremendous progress
- The recent success can be attributed to:
  - Explosion of data
  - Cheap computing cost – CPUs and GPUs
  - Improvement of machine learning models
- Much of the current excitement concerns a subfield of it called “deep learning”.

A brief history of Machine learning

- Most of the machine learning methods are based on supervised learning

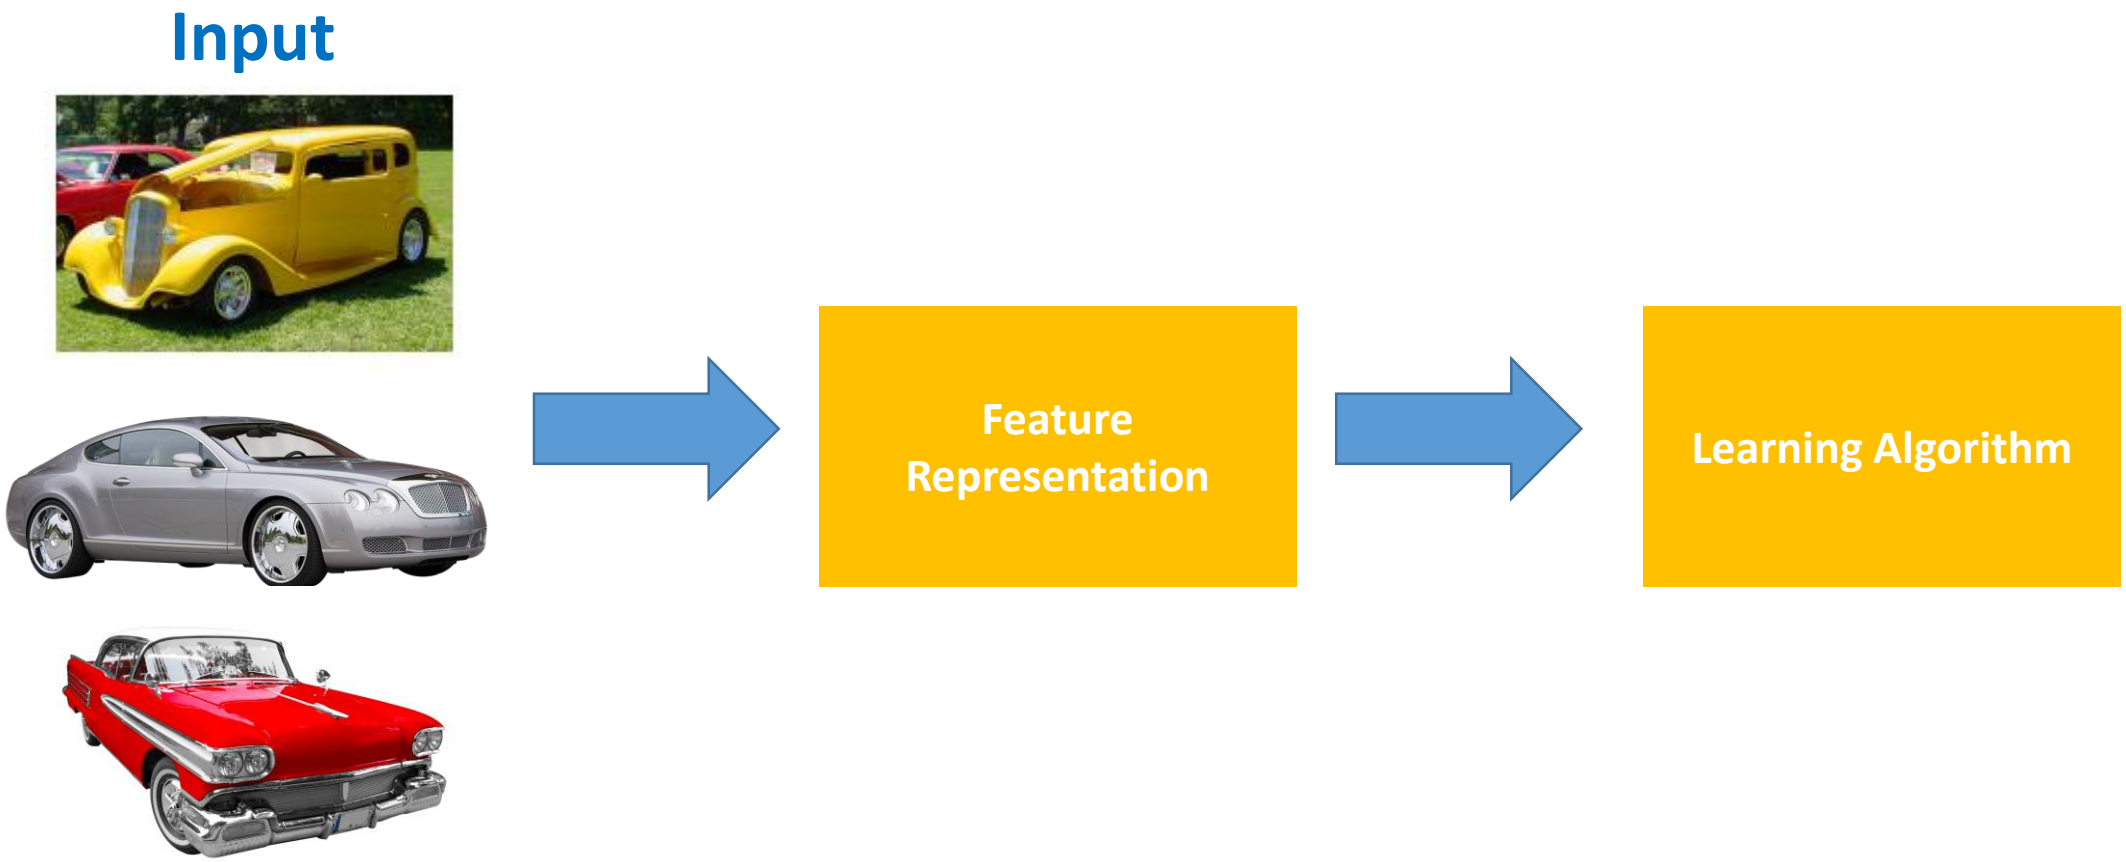

A brief history of Machine learning

- **The traditional model of pattern recognition (since the late 50's)**
  - Fixed/engineered features (or fixed kernel) + trainable classifier

hand-crafted Feature Extractor

"Simple" Trainable Classifier

32 45 21 ..
12 10 45 ..
17 33 36 ..
... ...

12 56 18 ..
92 76 22 ..
33 63 71 ..
... ...

Features Training

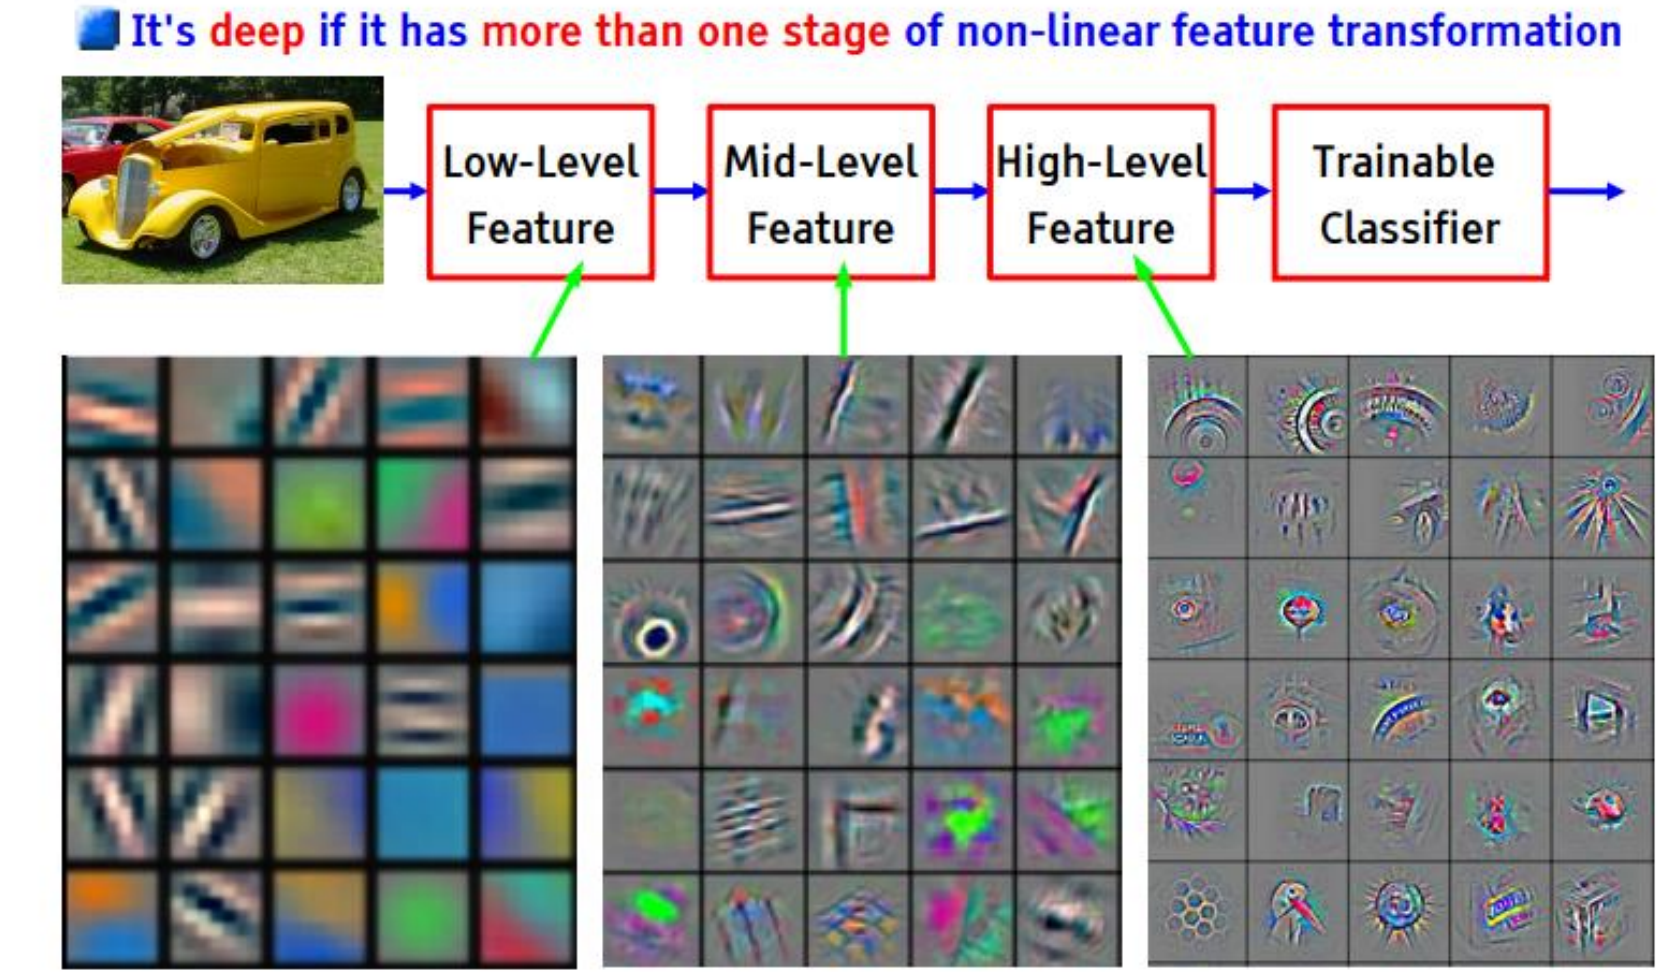

Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Trainable Feature Hierarchy

- Hierarchy of representations with increasing level of abstraction
- Each stage is a kind of trainable feature transform
- Image recognition
  - Pixel → edge → texton → motif → part → object
- Text
  - Character → word → word group → clause → sentence → story
- Speech
  - Sample → spectral band → sound → … → phone → phoneme → word

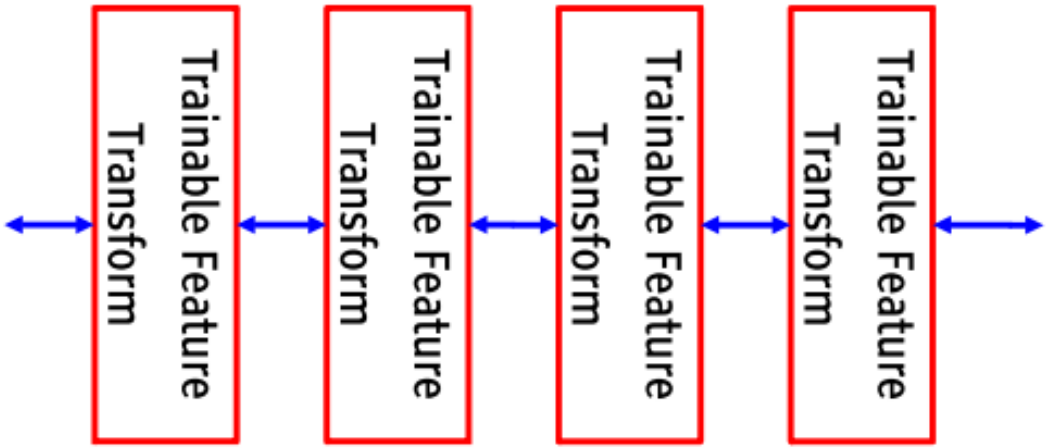

Traditional machine perception

- Hand crafted feature extractors

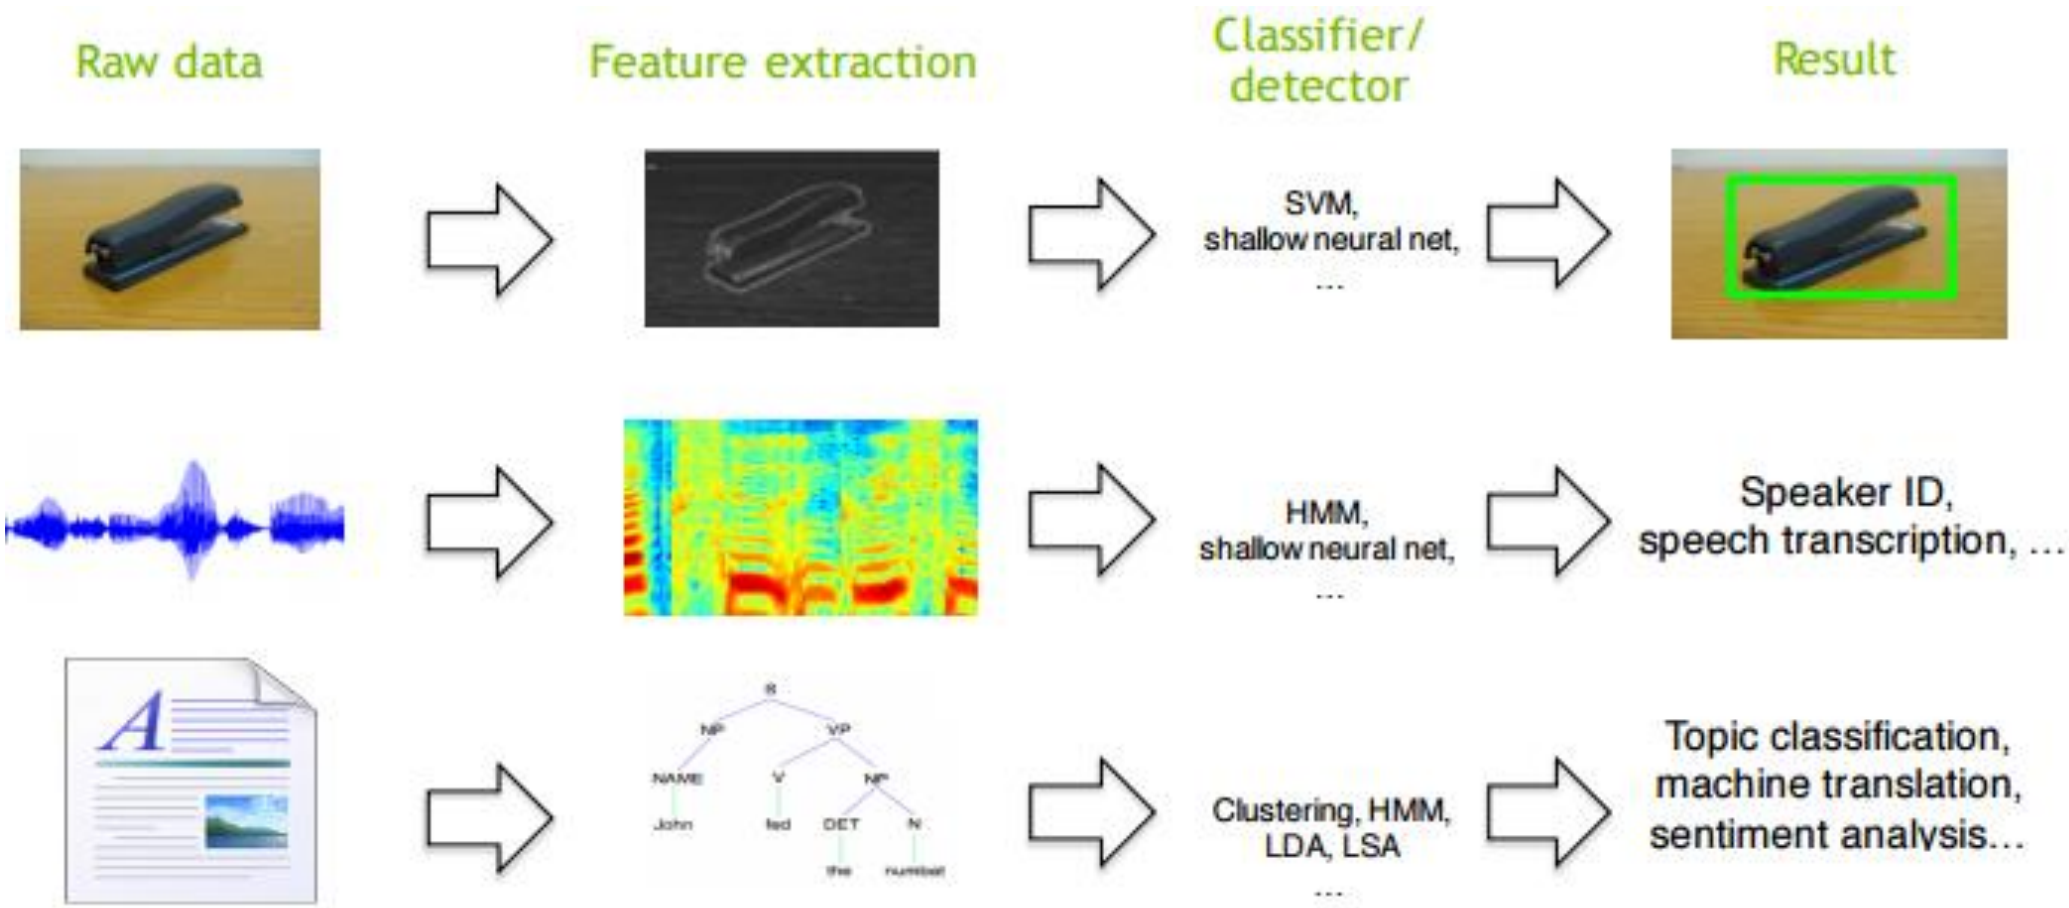

Human Brain

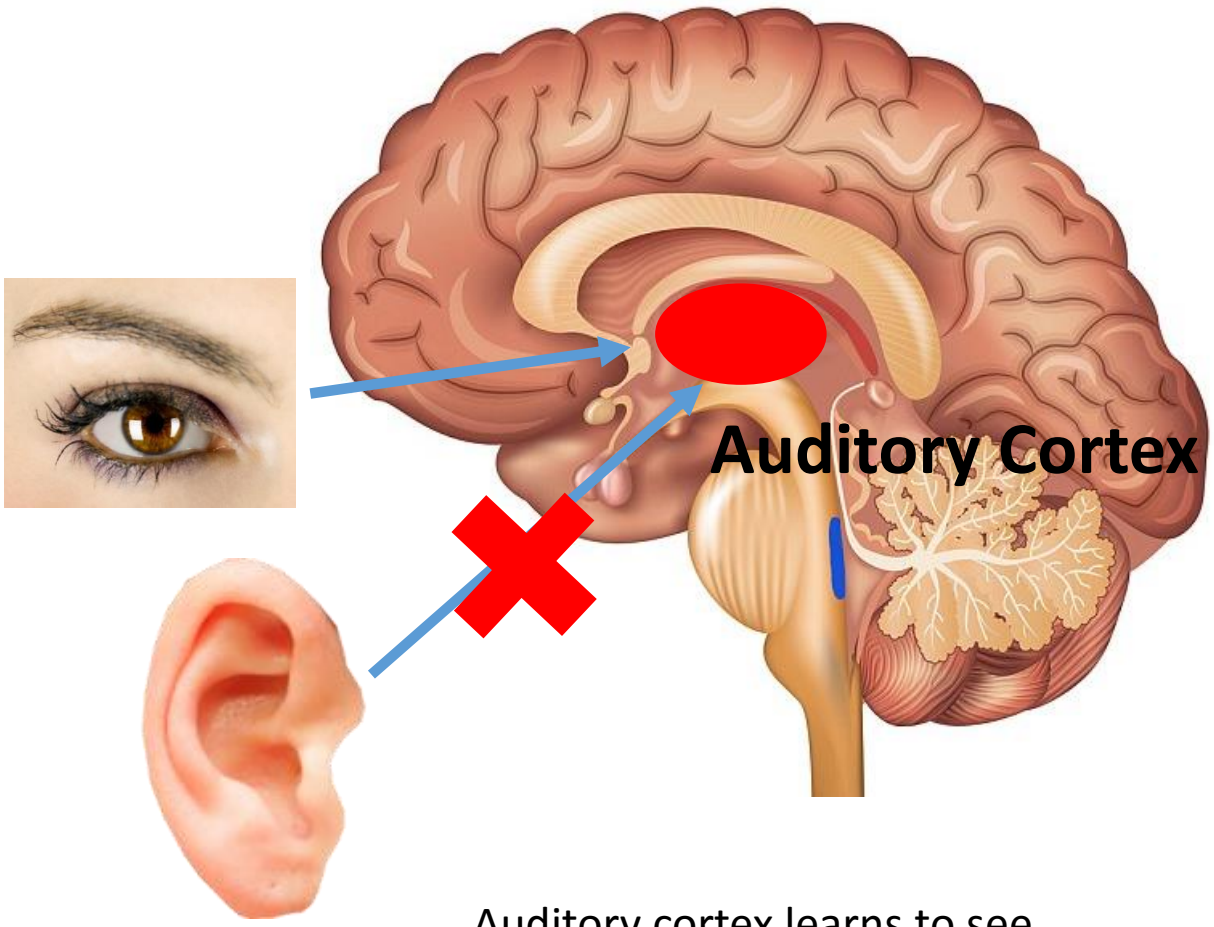

Auditory cortex learns to see. (Same rewiring process also works for touch/ somatosensory cortex.)

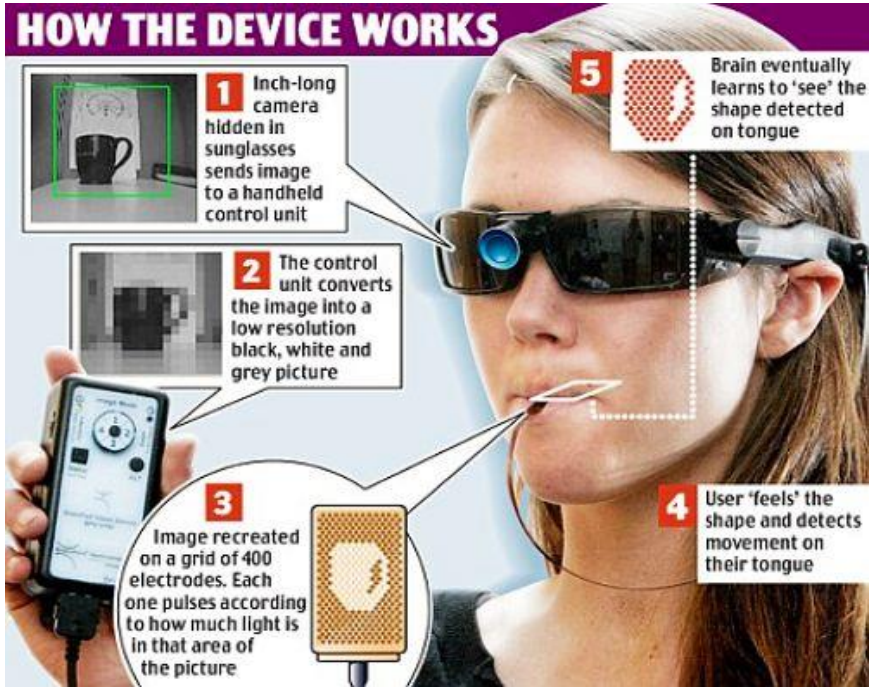

Seeing with tongue

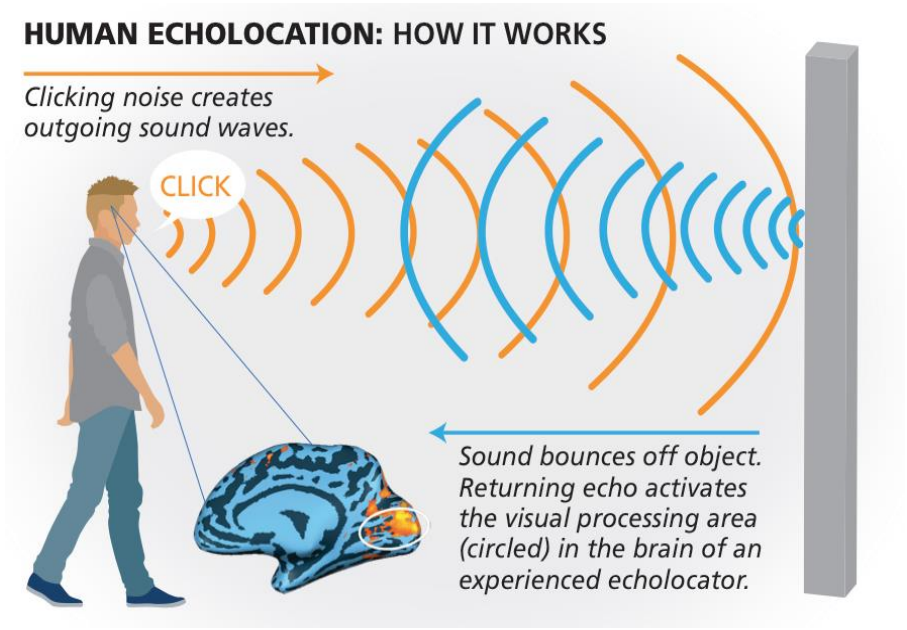

Human Brain

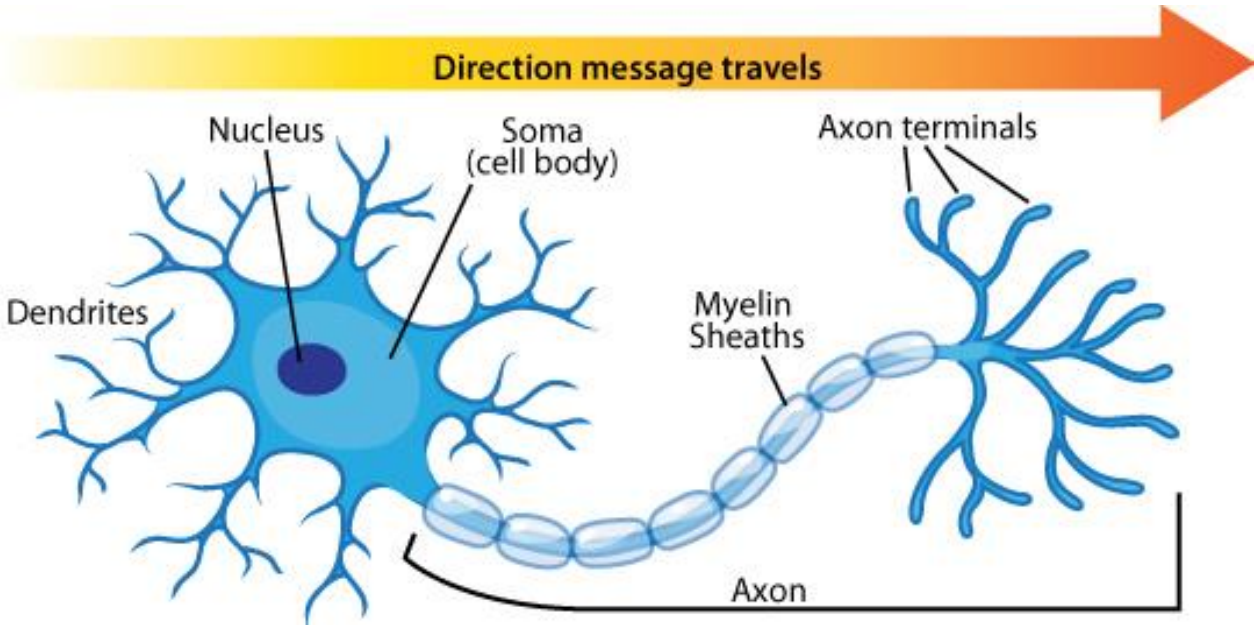

Biological Neuron

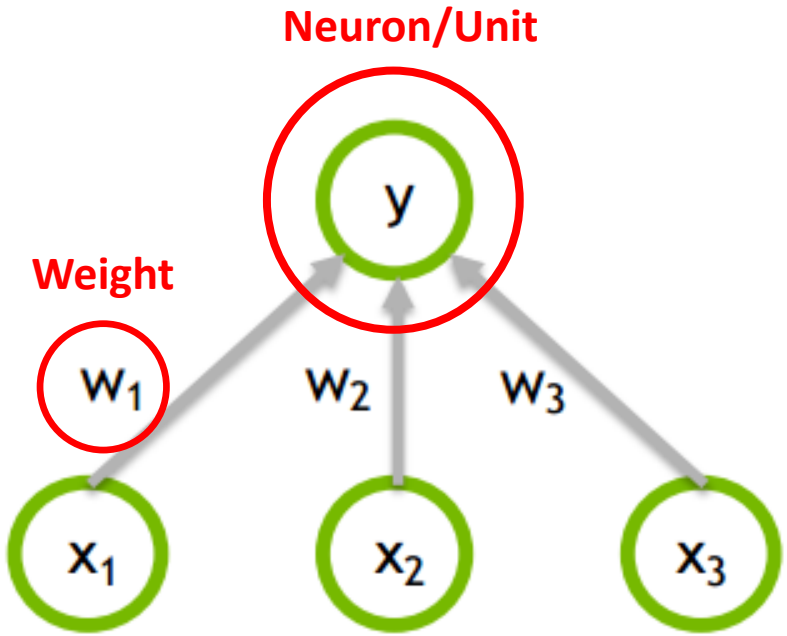

$$y=F(w_1x_1+w_2x_2+w_3x_3)$$

$$F(x)=\max(0,x)$$

Artificial Neuron

Neural Network

- Deep Learning is primarily about neural networks, where a network is an interconnected web of nodes and edges.
- Neural nets were designed to perform complex tasks, such as the task of placing objects into categories based on a few attributes.
- Neural nets are highly structured networks, and have three kinds of layers - an input, an output, and so called hidden layers, which refer to any layers between the input and the output layers.
- Each node (also called a neuron) in the hidden and output layers has a classifier.

Neural Network

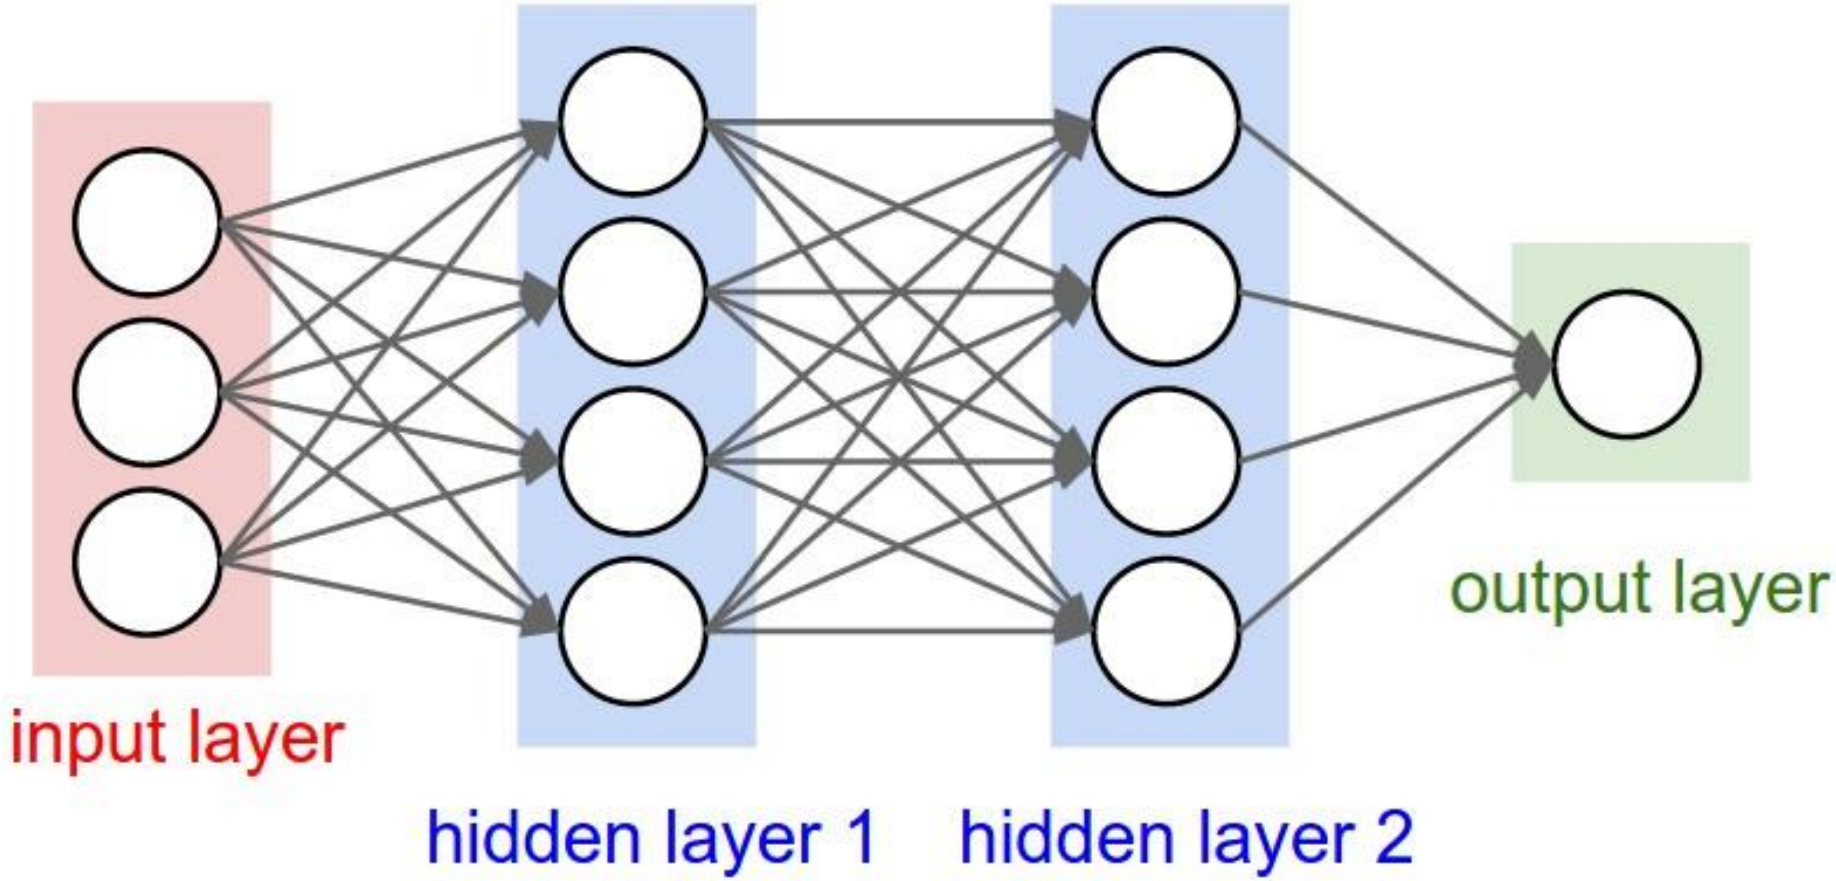

Neural Network: Forward Propagation

- The input neurons first receive the data features of the object. After processing the data, they send their output to the first hidden layer.
- The hidden layer processes this output and sends the results to the next hidden layer.
- This continues until the data reaches the final output layer, where the output value determines the object's classification.
- This entire process is known as Forward Propagation, or Forward prop.

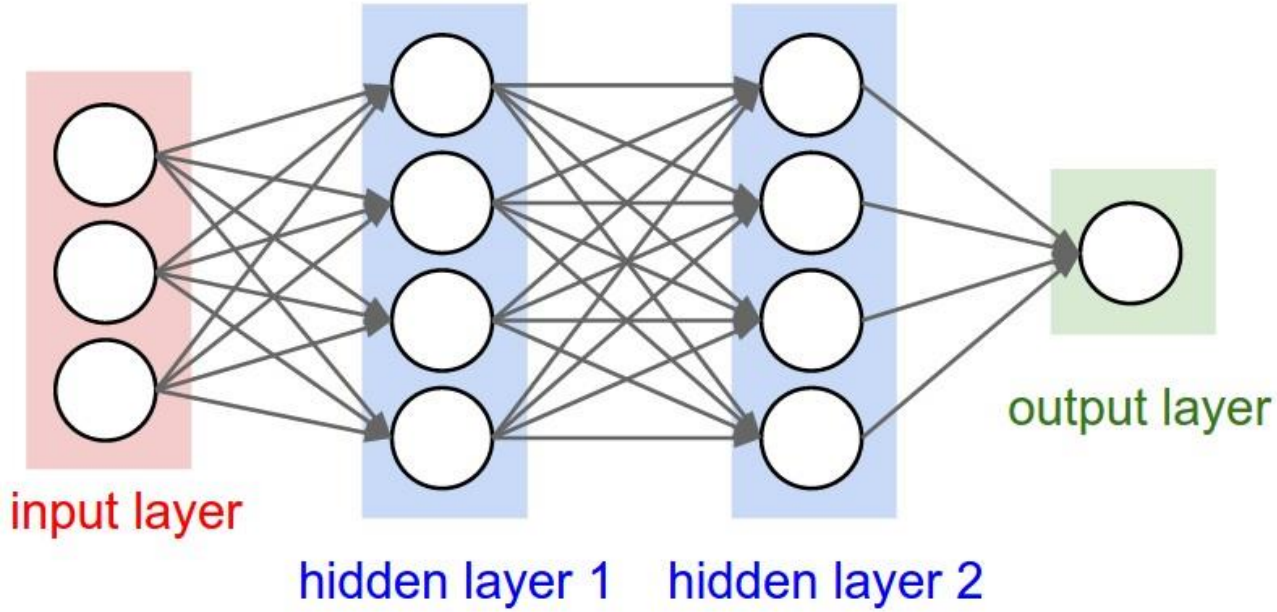

Neural Network: Backward Propagation

- To train a neural network over a large set of labelled data, you must continuously compute the difference between the network's predicted output and the actual output.
- This difference is called the cost, and the process for training a net is known as backpropagation, or backprop
- During backprop, weights and biases are tweaked slightly until the lowest possible cost is achieved.
- An important aspect of this process is the gradient, which is a measure of how much the cost changes with respect to a change in a weight or bias value.

The 1990s view of what was wrong with back-propagation

- It required a lot of labelled training data
  - almost all data is unlabeled
- The learning time did not scale well
  - It was very slow in networks with multiple hidden layers.
- It got stuck at local optima
  - These were often surprisingly good but there was no good theory

Deep Belief Network (DBN)

- The Deep Belief Network, or DBN, was also conceived by Geoff Hinton.
- Used by Google for their work on the image recognition problem.
- DBN is trained two layers at a time, and these two layers are treated like an RBM.
- Throughout the net, the hidden layer of an RBM acts as the input layer of the adjacent one. So the first RBM is trained, and its outputs are then used as inputs to the next RBM. This procedure is repeated until the output layer is reached.

DBNs are stacks of restricted Boltzmann machines forming deep (multi-layer) architecture.

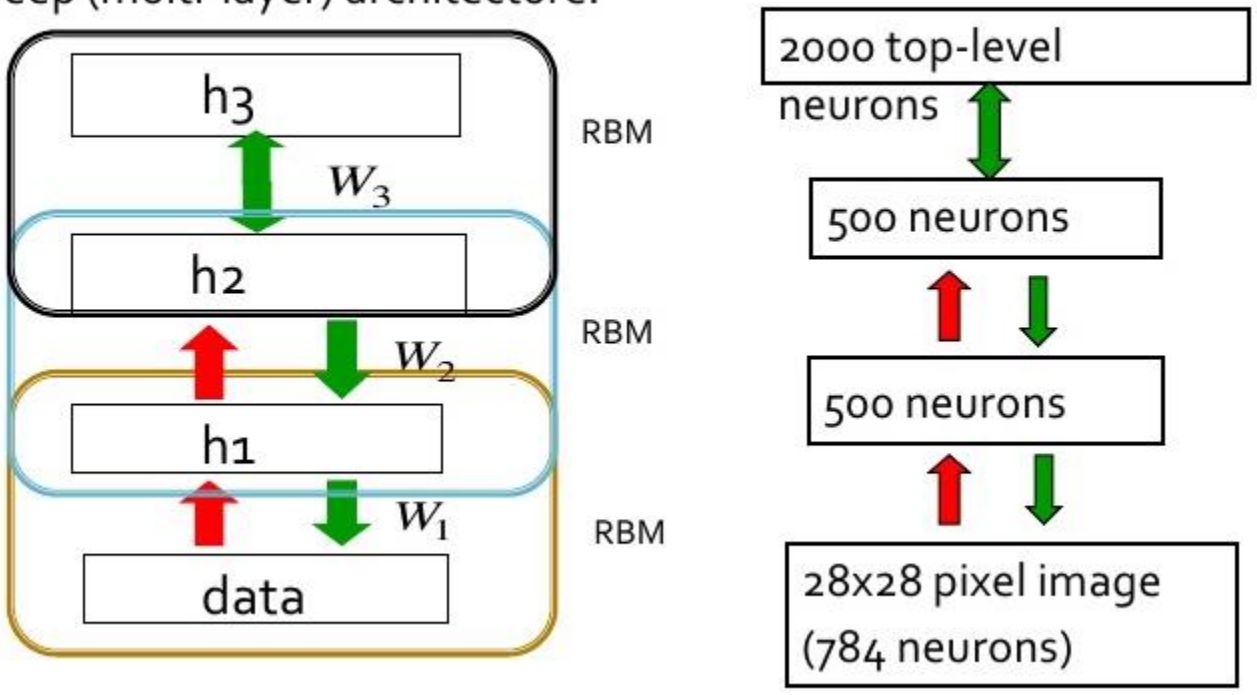

Deep Belief Network (DBN)

- DBN is capable of recognizing the inherent patterns in the data. In other words, it’s a sophisticated, multilayer feature extractor.
- The unique aspect of this type of net is that each layer ends up learning the full input structure.
- Layers generally learn progressively complex patterns – for facial recognition, early layers could detect edges and later layers would combine them to form facial features.
- DBN learns the hidden patterns globally, like a camera slowly bringing an image into focus.
- DBN still requires a set of labels to apply to the resulting patterns. As a final step, the DBN is fine-tuned with supervised learning and a small set of labeled examples.

Deep Neural Network (Deep Net)

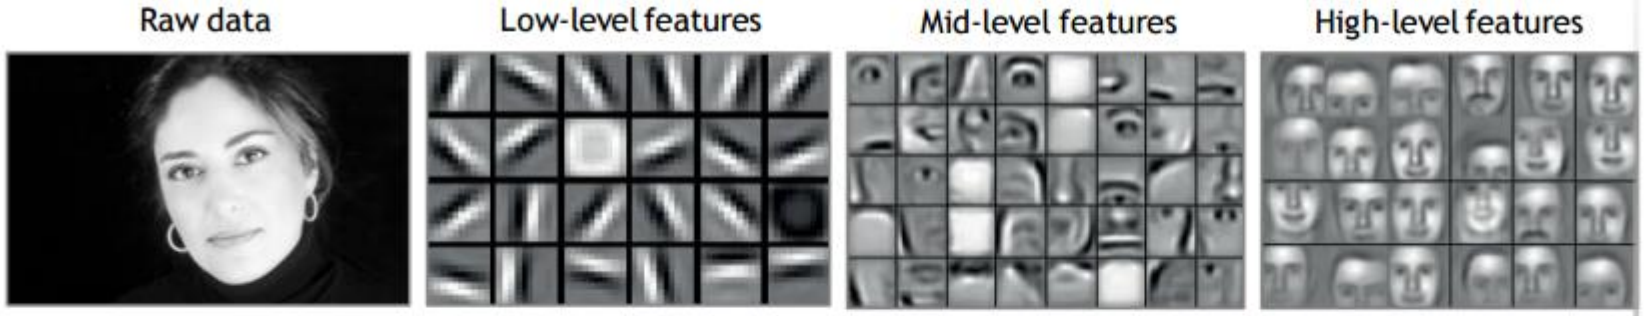

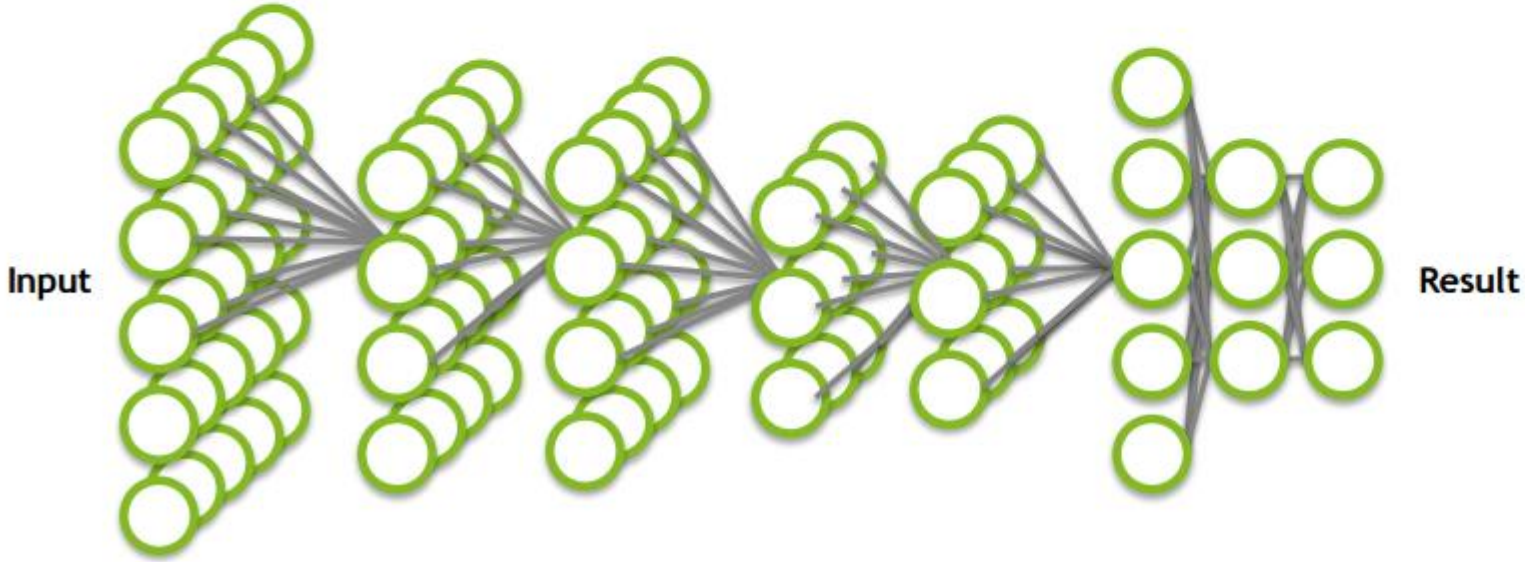

Application components:

Task objective
- e.g. Identify face

Training data
- 10-100M images

Network architecture
- ~10 layers
- 1B parameters

Learning algorithm
- ~30 Exaflops
- ~30 GPU days

Convolutional Neural Network (CNN)

- CNN inspired by the Visual Cortex.
- CNNs are deep nets that are used for image, object, and even speech recognition.
- Pioneered by Yann Lecun (NYU)
- Deep supervised neural networks are generally too difficult to train.
- CNNs have multiple types of layers, the first of which is the convolutional layer.

Convolutional Neural Network (CNN)

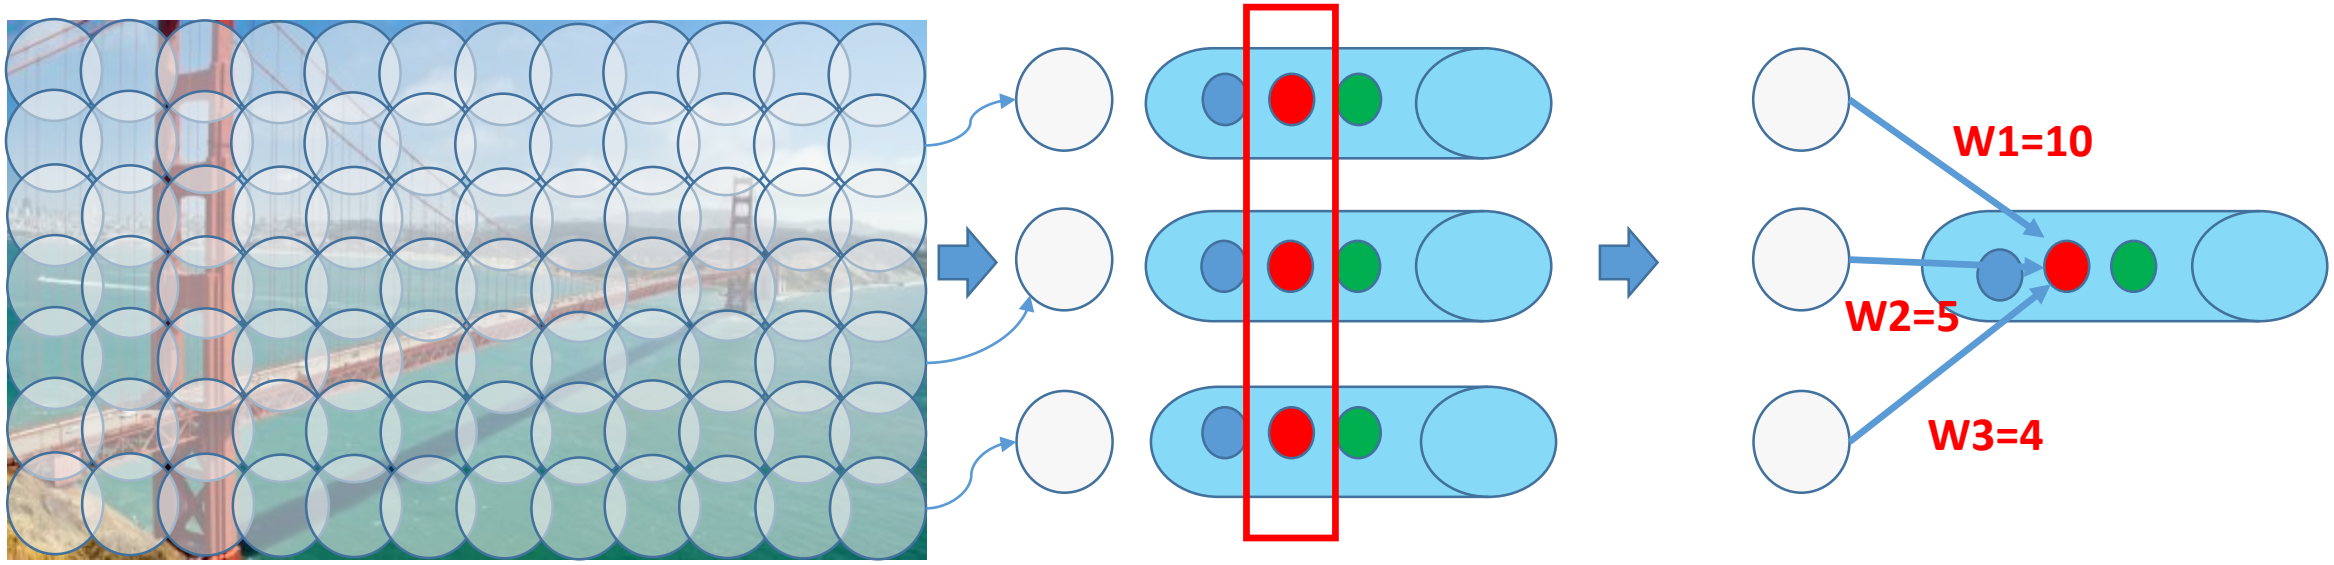

Filter 2/ Neural 2

- A series of filters forms layer one, called the convolutional layer. The weights and biases in this layer determine the effectiveness of the filtering process.
- Each flashlight represents a single neuron. Typically, neurons in a layer activate or fire. On the other hand, in the convolutional layer, neurons search for patterns through convolution. Neurons from different filters search for different patterns, and thus they will process the input differently.

Convolutional Neural Network (CNN)

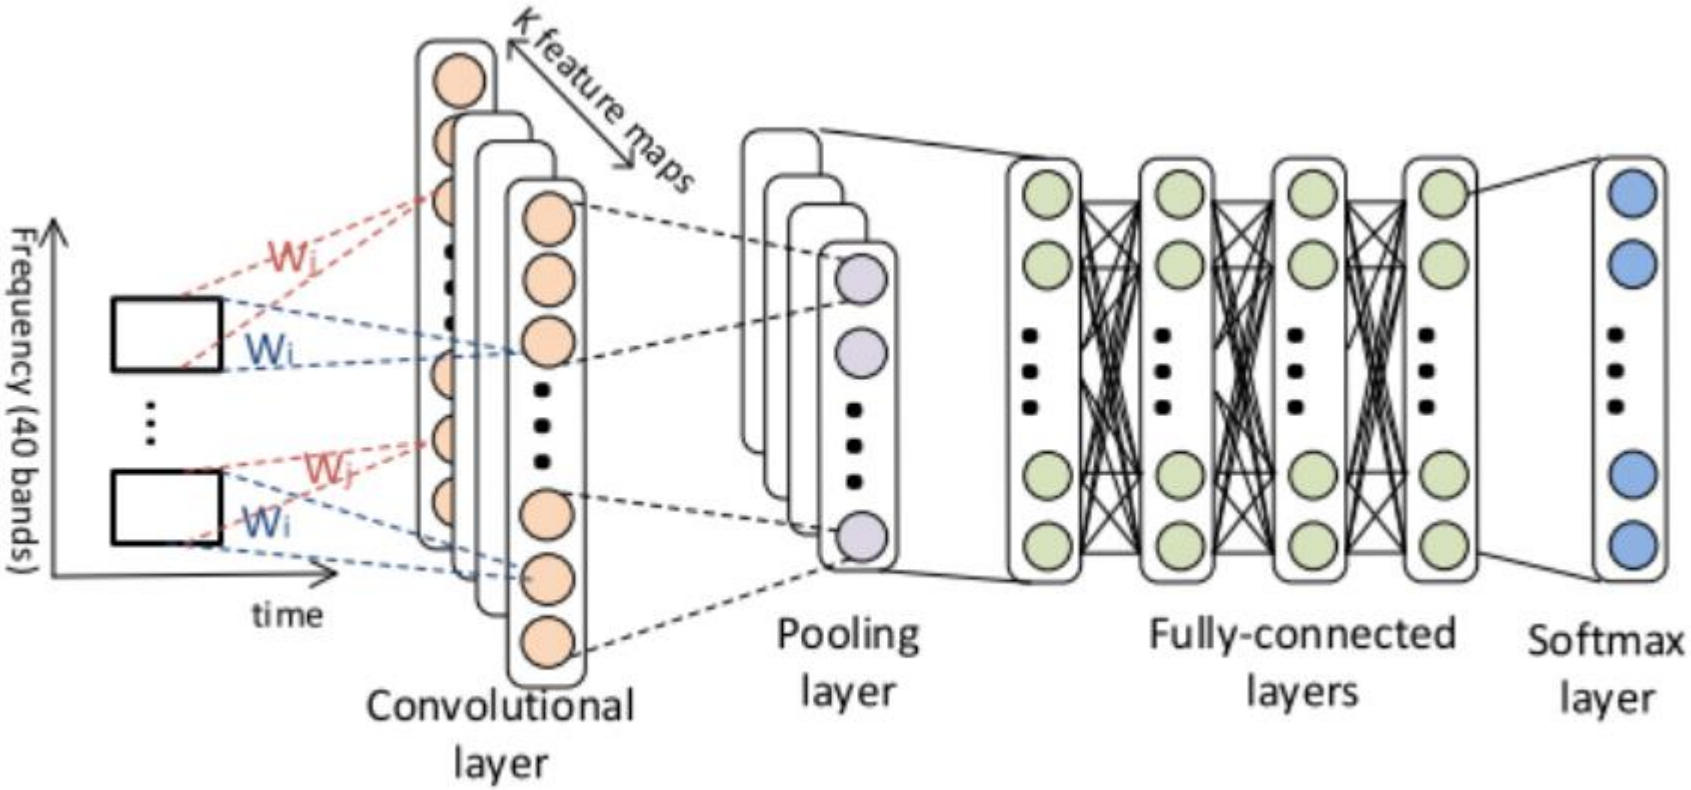

CNN: Application

- Classify a scene in an image
  - Image Classifier Demo (NYU): http://horatio.cs.nyu.edu/
- Describe or understanding an image
  - Toronto Deep Learning Demo: http://deeplearning.cs.toronto.edu/i2t
  - MIT Scene Recognition Demo: http://places.csail.mit.edu/demo.html
- Handwriting recognition
  - Handwritten digits recognition: http://cs.stanford.edu/people/karpathy/convnetjs/demo/mnist.html
- Video classification
  - Large-scale Video Classification with Convolutional Neural Networks http://cs.stanford.edu/people/karpathy/deepvideo/

Recurrent Neural Network (RNN)

- The Recurrent Neural Net (RNN) is the brainchild of Juergen Schmidhuber and Sepp Hochreiter.
- RNNs have a feedback loop where the net's output is fed back into the net along with the next input.
- RNNs receive an input and produce an output. Unlike other nets, the inputs and outputs can come in a sequence.
- Variant of RNN is Long Term Short Memory (LSTM)

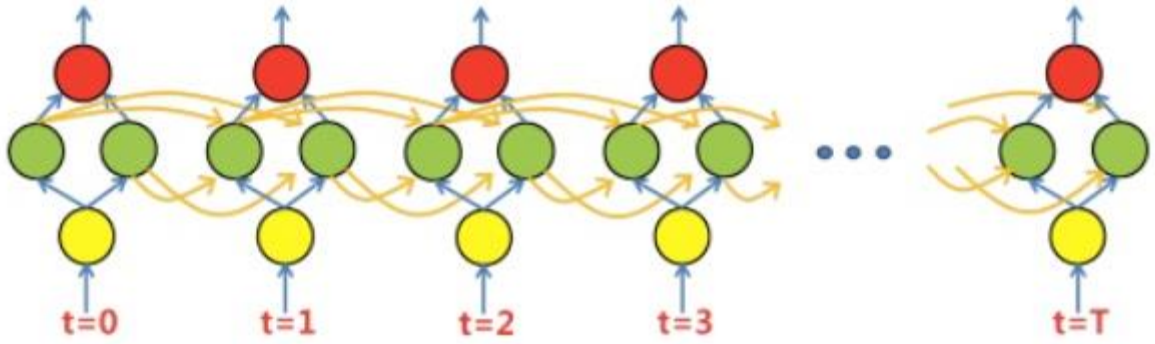

RNN: Application

- RNN is suitable for time series data, where an output can be the next value in a sequence, or the next several values

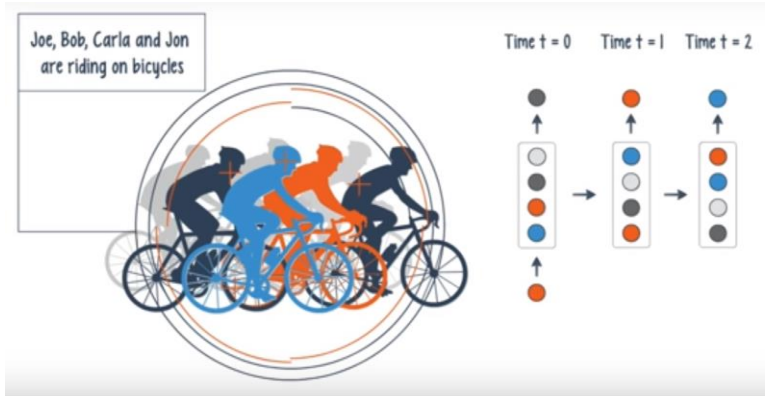

Image captioning

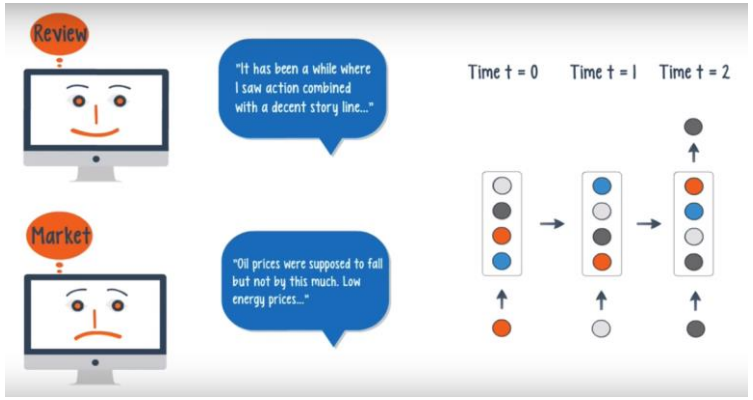

Document Classification

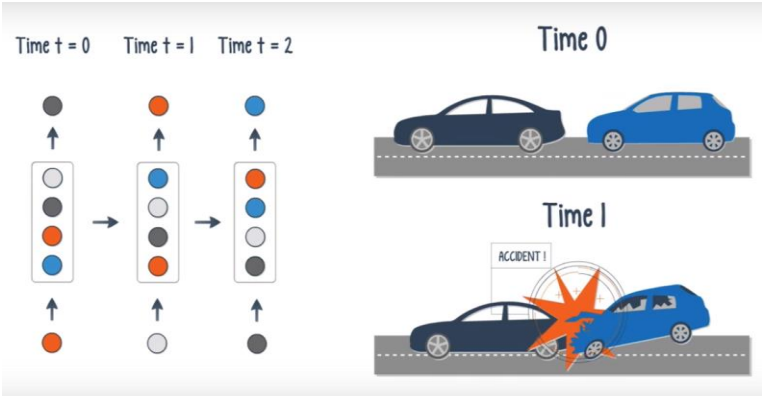

Classify Image frame by frame

Deep Learning: Benefits

- Robust
  - No need to design the features ahead of time – features are automatically learned to be optimal for the task at hand
  - Robustness to natural variations in the data is automatically learned
- Generalizable
  - The same neural net approach can be used for many different applications and data types
- Scalable
  - Performance improves with more data, method is massively parallelizable

Deep Learning: Weaknesses

- Deep Learning requires a large dataset, hence long training period.
- In term of cost, Machine Learning methods like SVMs and other tree ensembles are very easily deployed even by relative machine learning novices and can usually get you reasonably good results.
- Deep learning methods tend to learn everything. It's better to encode prior knowledge about structure of images (or audio or text).
- The learned features are often difficult to understand. Many vision features are also not really human-understandable (e.g, concatenations/combinations of different features).
- Requires a good understanding of how to model multiple modalities with traditional tools.

Deep Learning: Applications

DEEP LEARNING EVERYWHERE

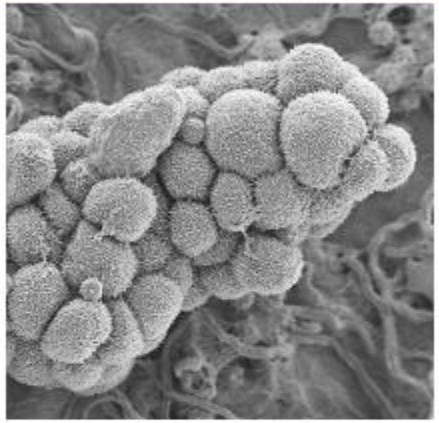

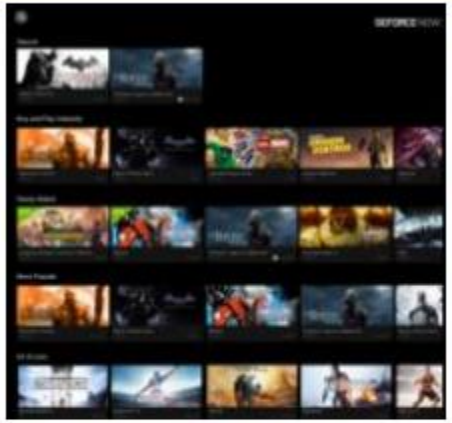

INTERNET & CLOUD

Image Classification
Speech Recognition
Language Translation
Language Processing
Sentiment Analysis
Recommendation

MEDICINE & BIOLOGY

Cancer Cell Detection
Diabetic Grading
Drug Discovery

MEDIA & ENTERTAINMENT

Video Captioning
Video Search
Real Time Translation

SECURITY & DEFENSE

Face Detection
Video Surveillance
Satellite Imagery

AUTONOMOUS MACHINES

Pedestrian Detection
Lane Tracking
Recognize Traffic Sign

Deep Learning: Application

https://deepmind.com/alpha-go

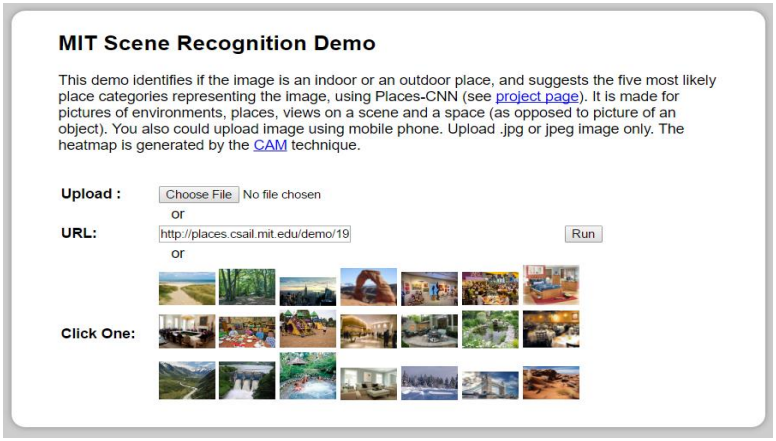

http://places.csail.mit.edu/demo.html

Pedestrian detection using DL

Robotic grasping

Deep Learning Development Cycle

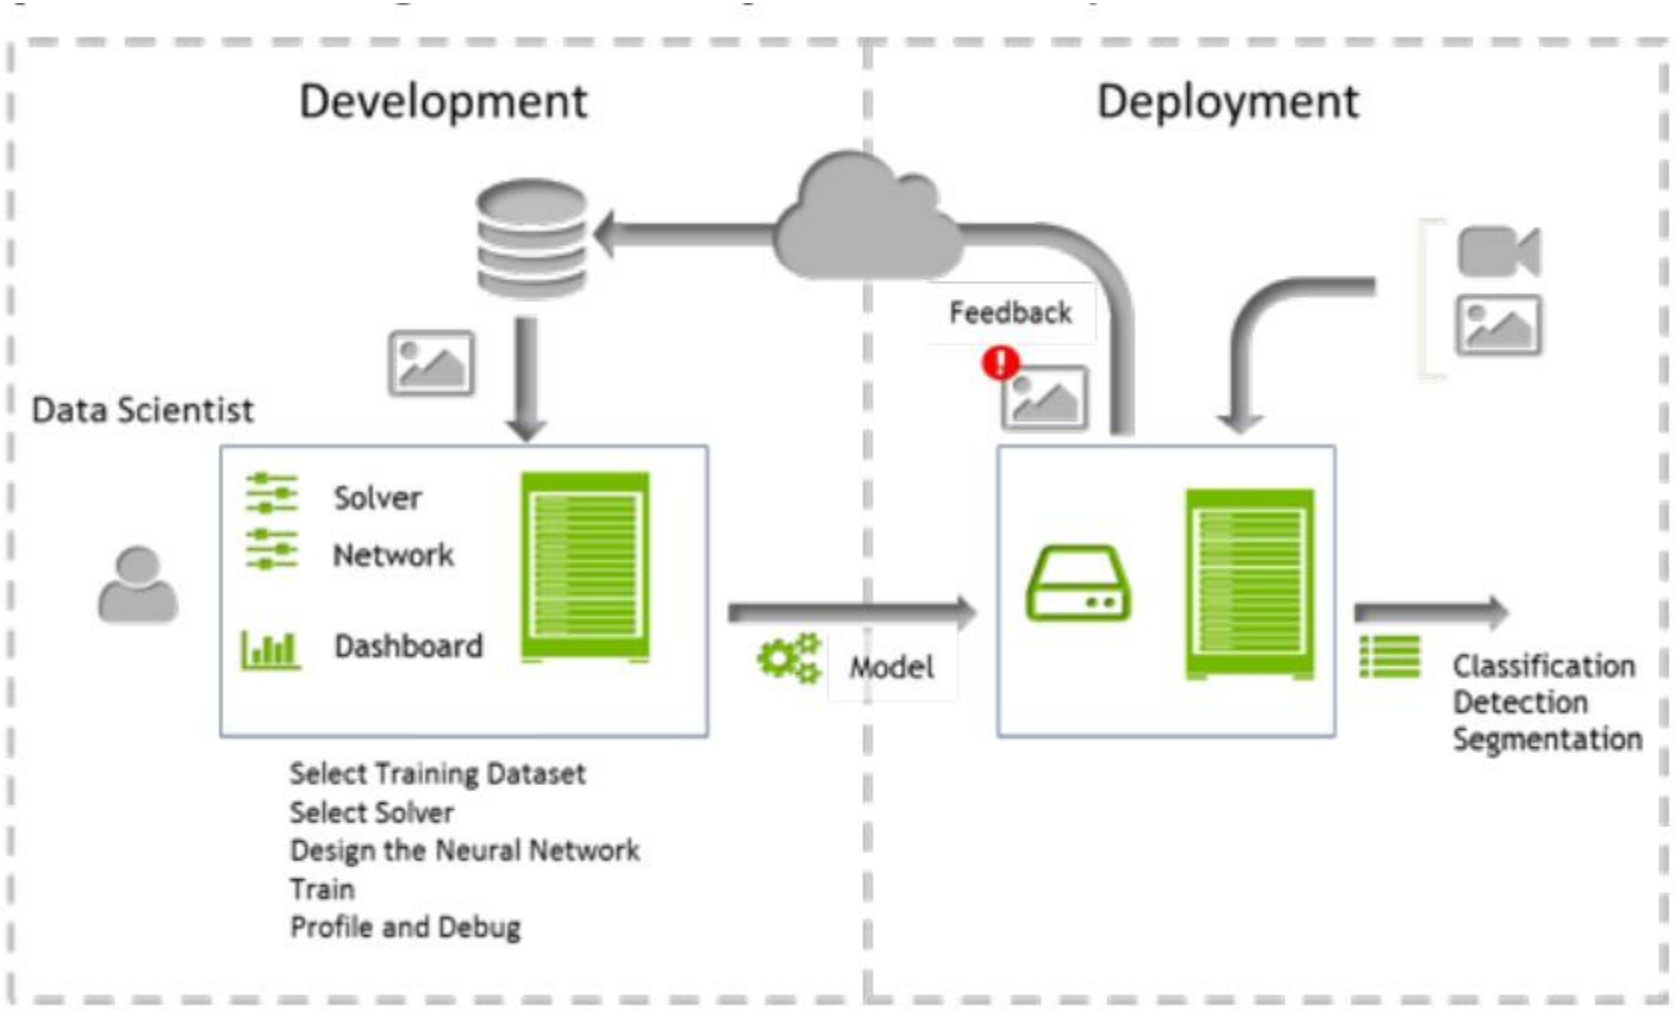

Deep Learning: Platform & Frameworks & Libraries

Platform

- Ersatz Labs - cloud-based deep learning platform [http://www.ersatz1.com/]
- H20 – deep learning framework that comes with R and Python interfaces [http://www.h2o.ai/verticals/algos/deep-learning/]

Framework

- Caffe - deep learning framework made with expression, speed, and modularity in mind. Developed by the Berkeley Vision and Learning Center (BVLC) [http://caffe.berkeleyvision.org/]
- Torch - scientific computing framework with wide support for machine learning algorithms that puts GPUs first. Based on Lua programming language [http://torch.ch/]

Library

- Tensorflow - open source software library for numerical computation using data flow graphs from Google [https://www.tensorflow.org/]
- Theano - a python library developed by Yoshua Bengio's team [http://deeplearning.net/software/theano/]

Learned Models

- Trained Models can be shared with others
- Save the training time
- For example: AlexNet, GoogLeNet, ParseNet, etc
- URLs:
  - https://github.com/BVLC/caffe/wiki/Model-Zoo
  - http://deeplearning4j.org/model-zoo

DEMO

Nvidia: Digits

- The NVIDIA Deep Learning GPU Training System (DIGITS) puts the power of deep learning in the hands of data scientists and researchers.
- Quickly design the best deep neural network (DNN) for your data using real-time network behavior visualization.
- https://developer.nvidia.com/digits

Car Park Images

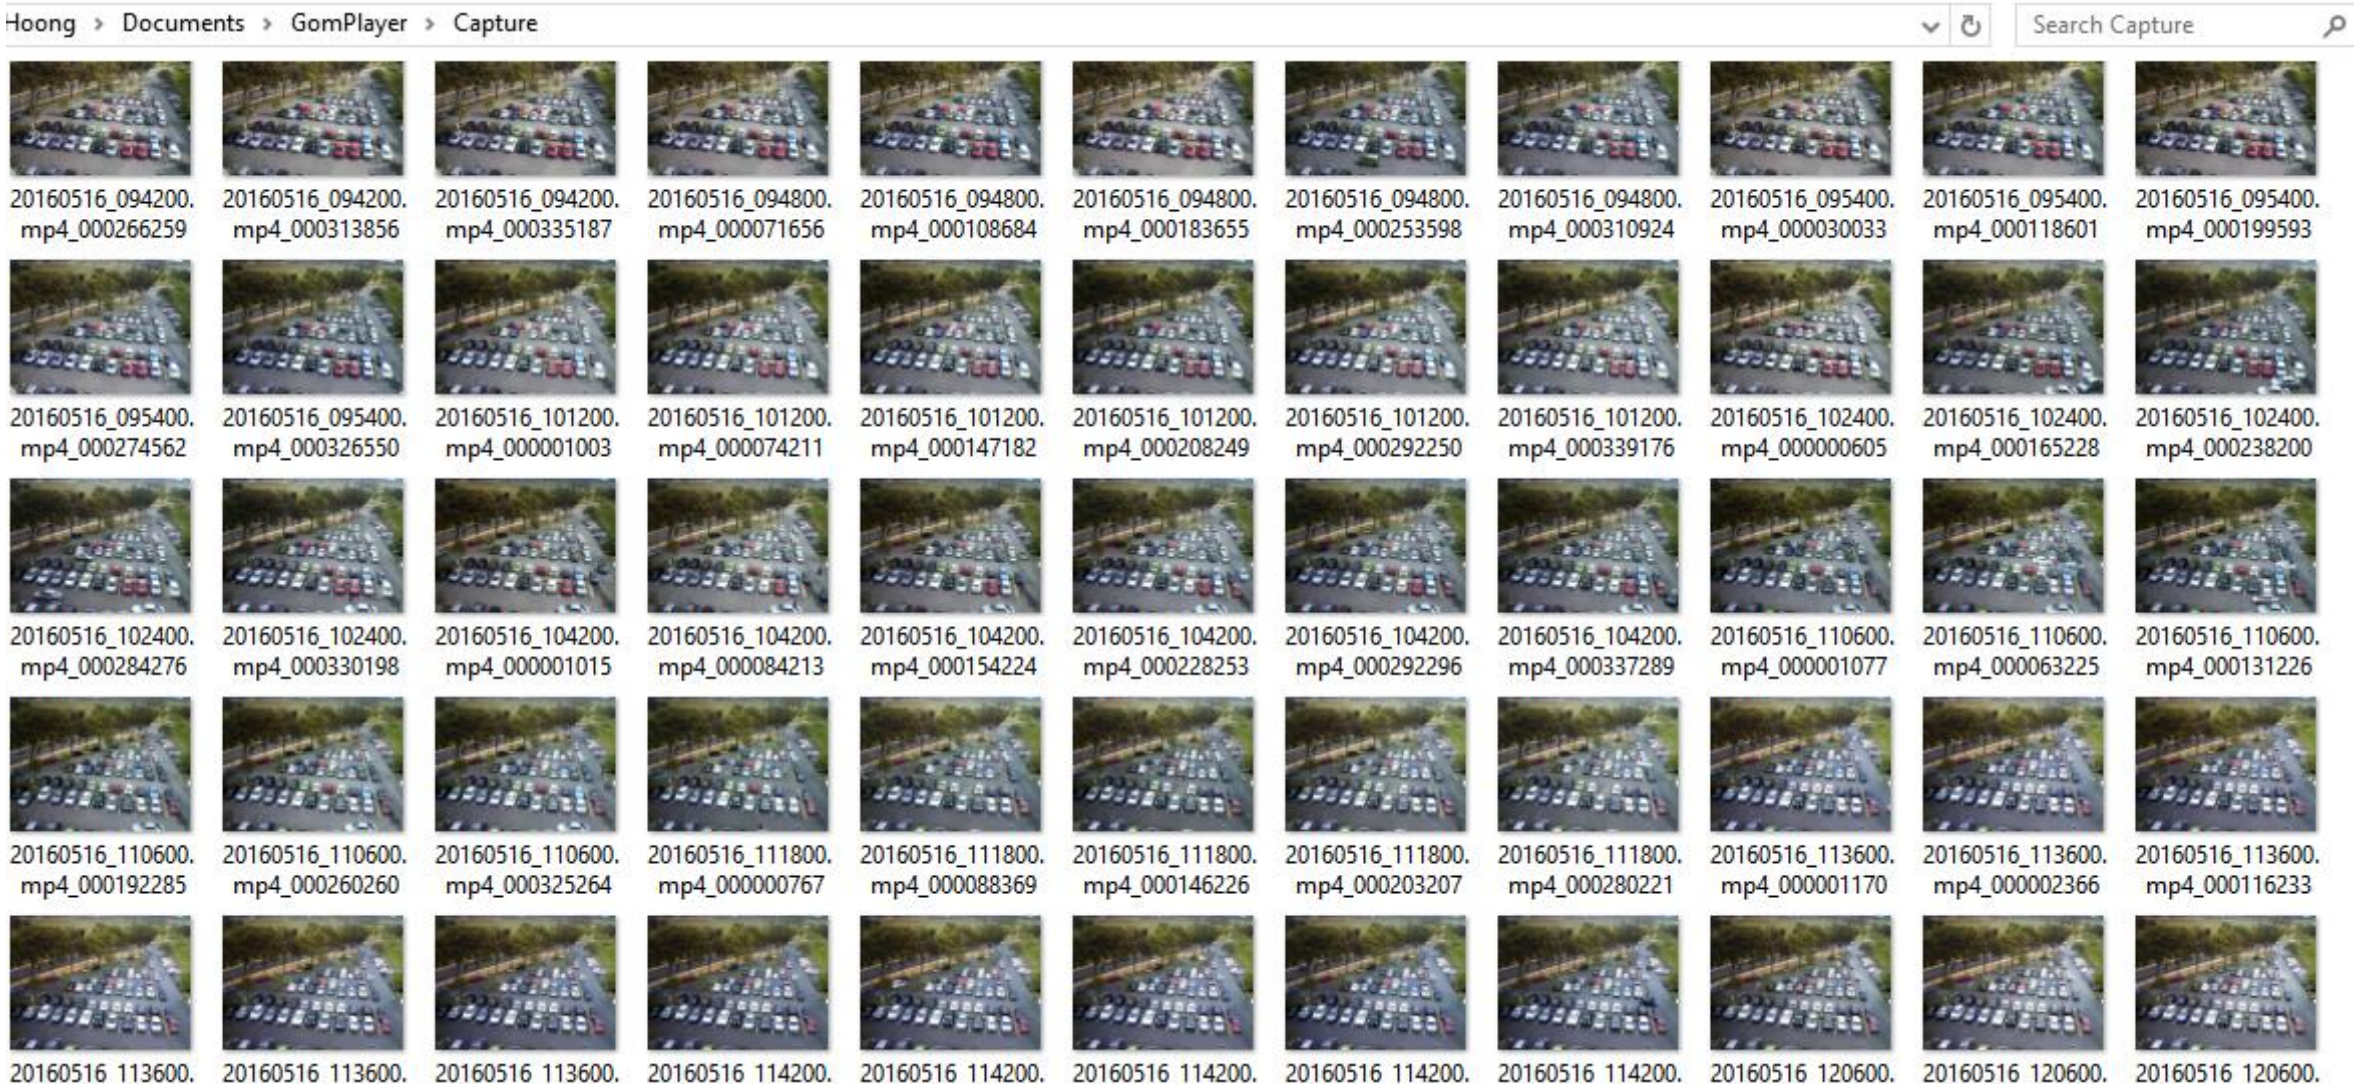

Car park images

Cropped Car Park space

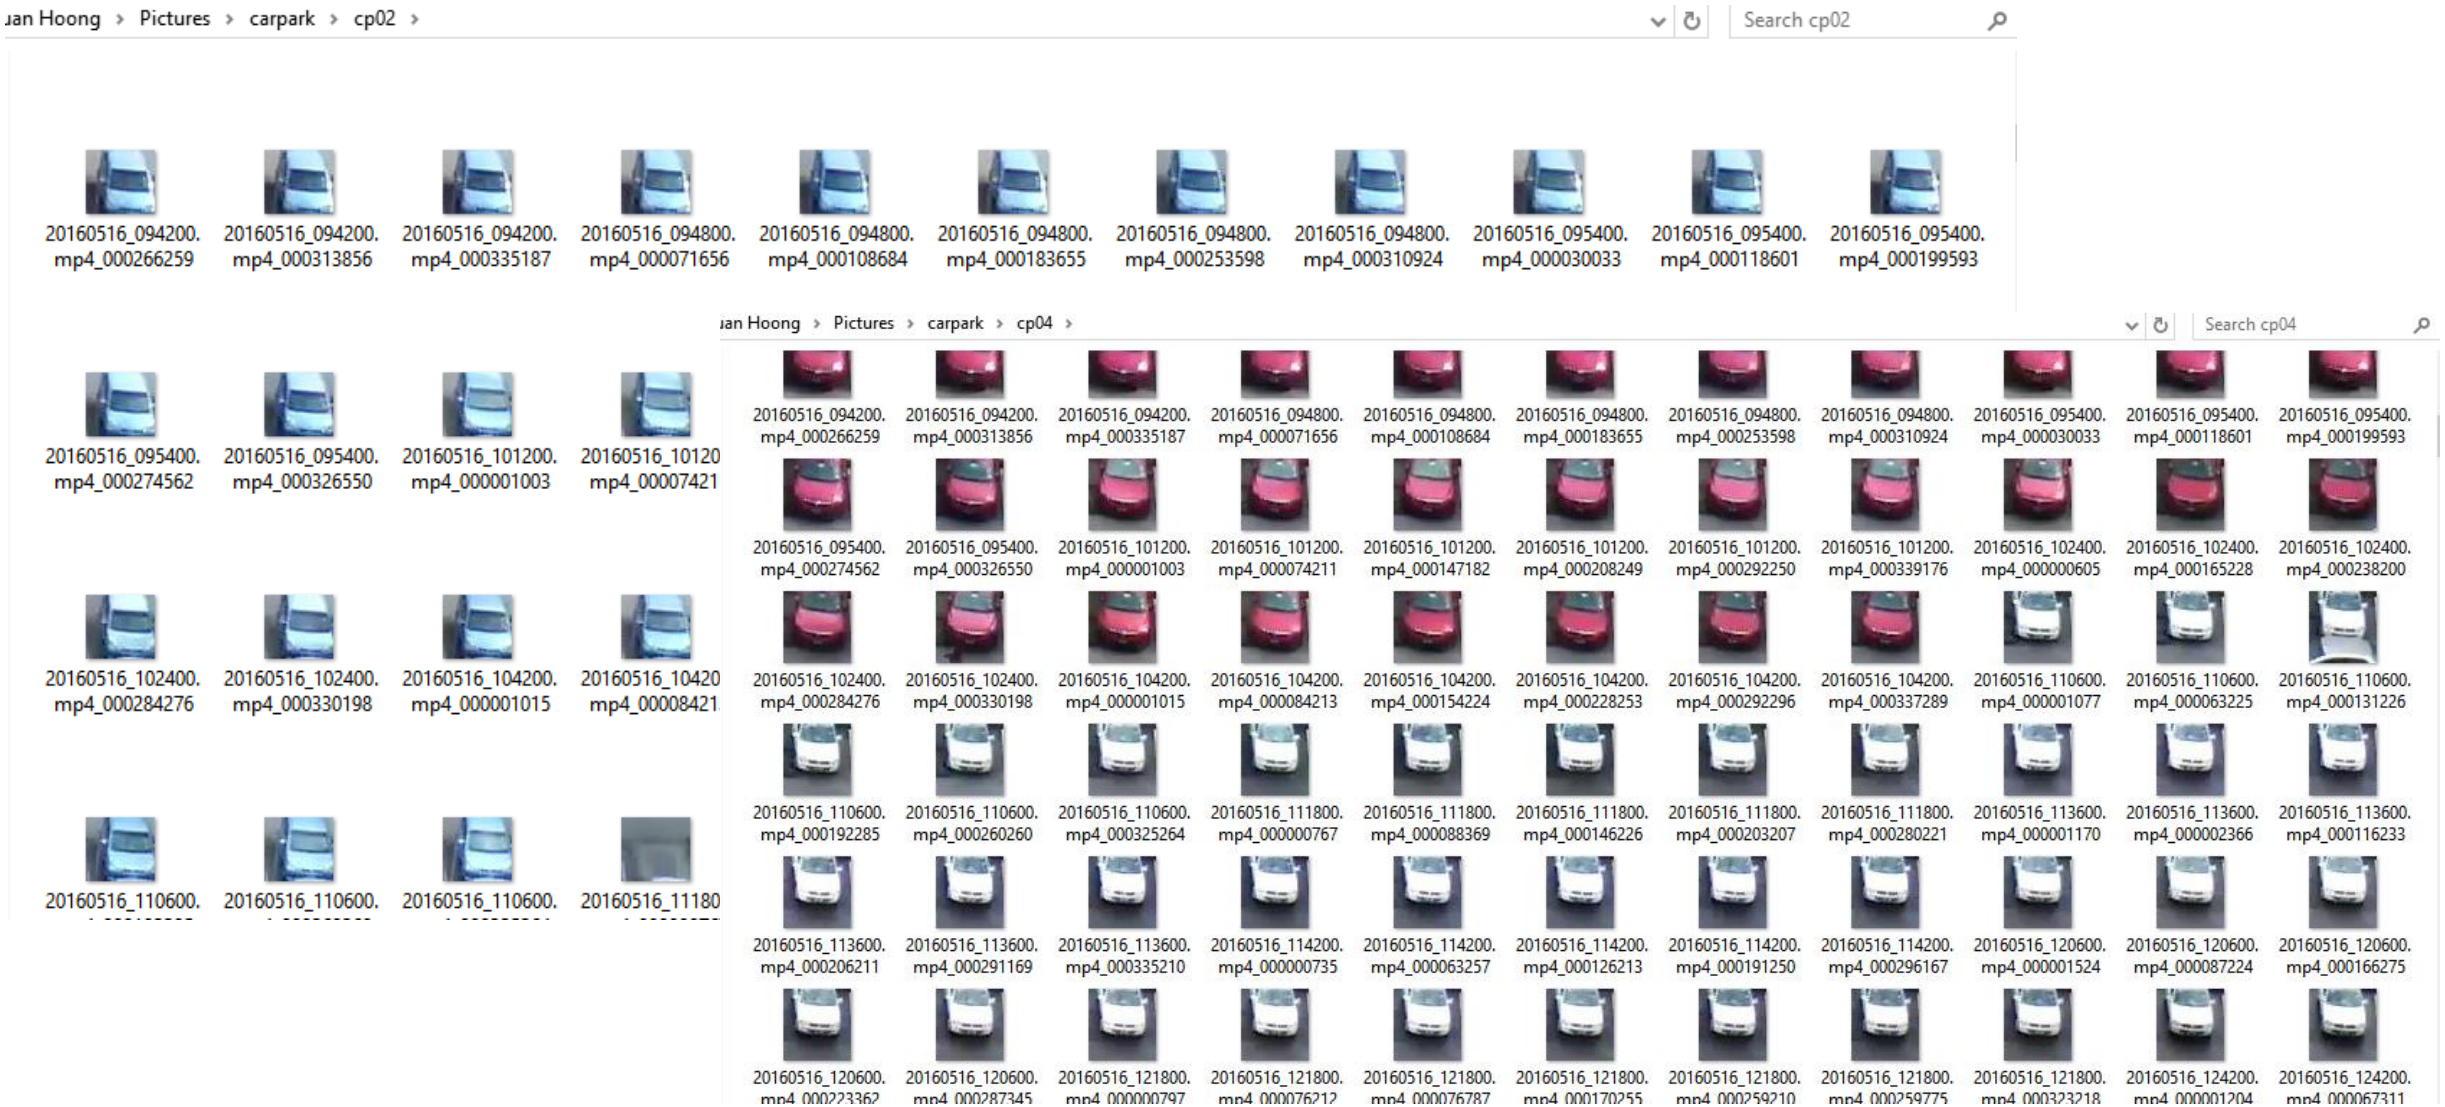

Digits – Image Classification Model

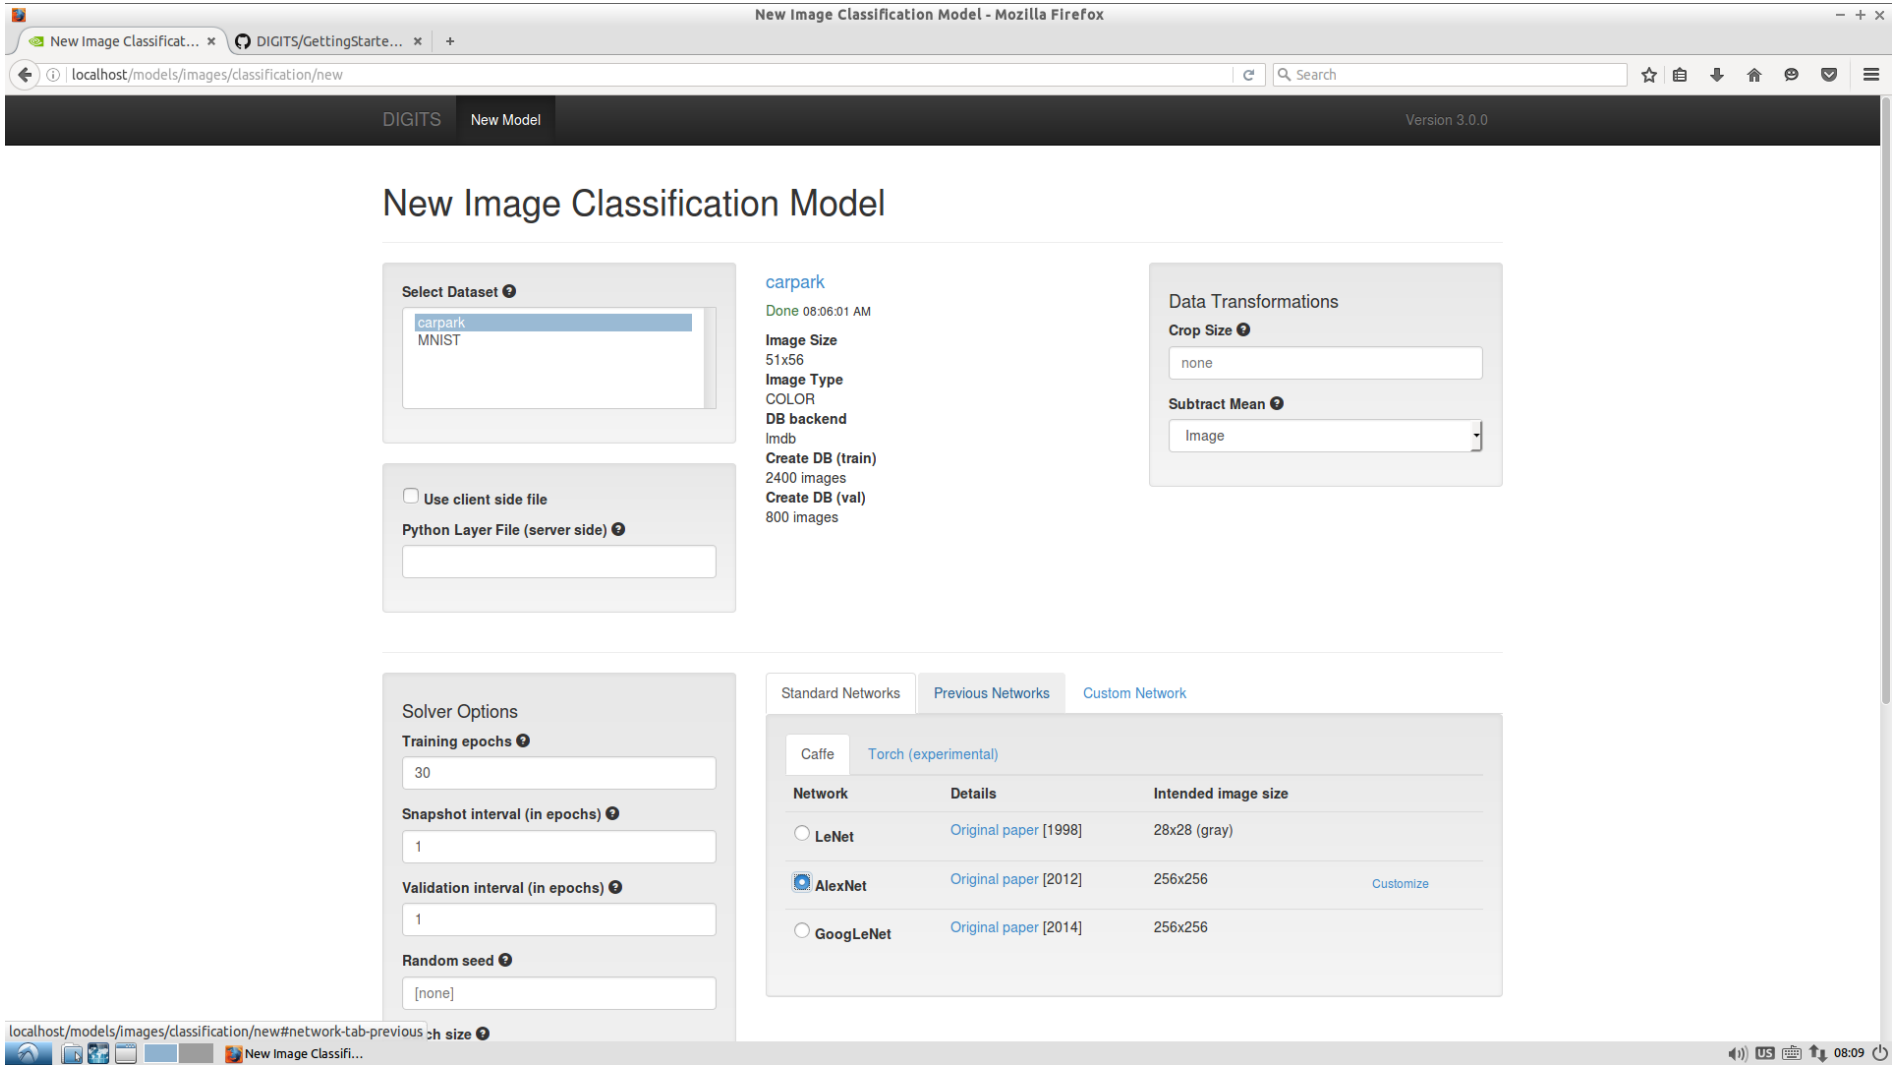

Digits – AlexNet Training

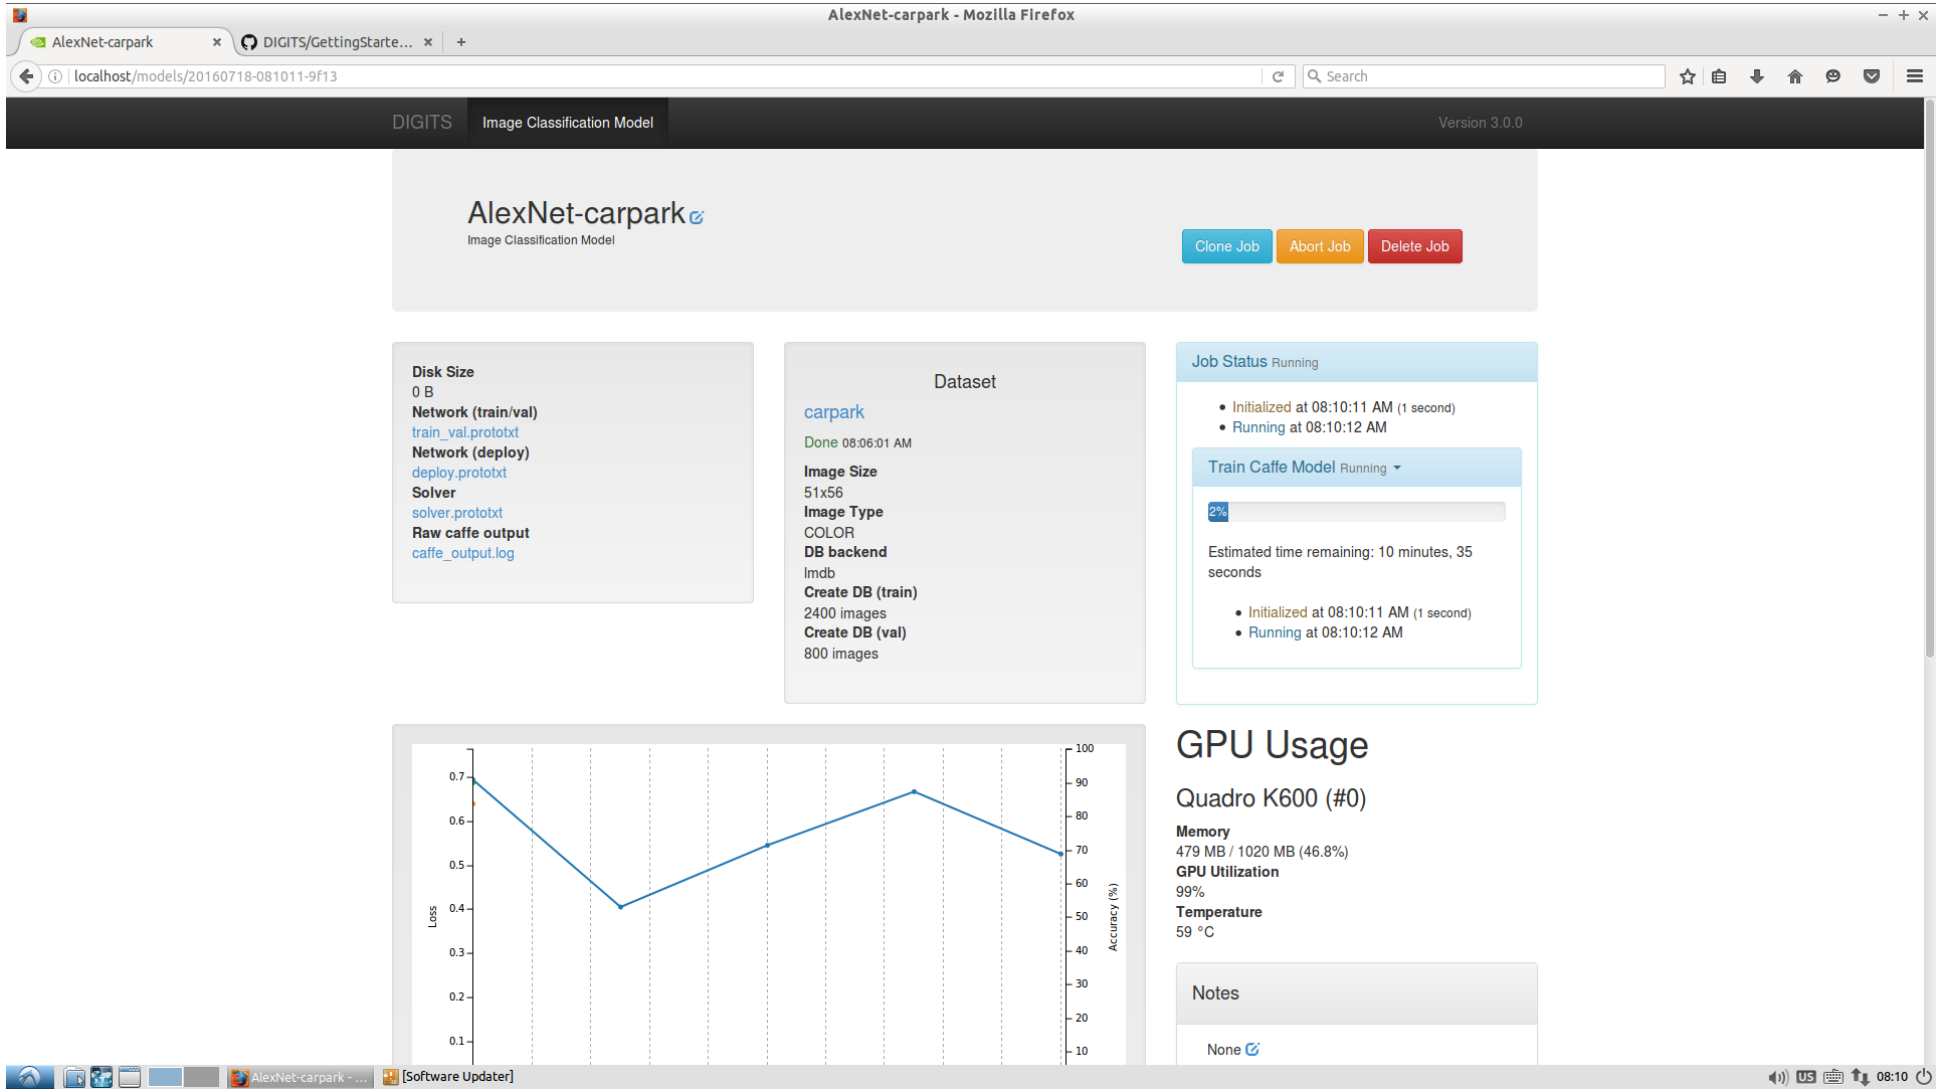

Digits – AlexNet Training

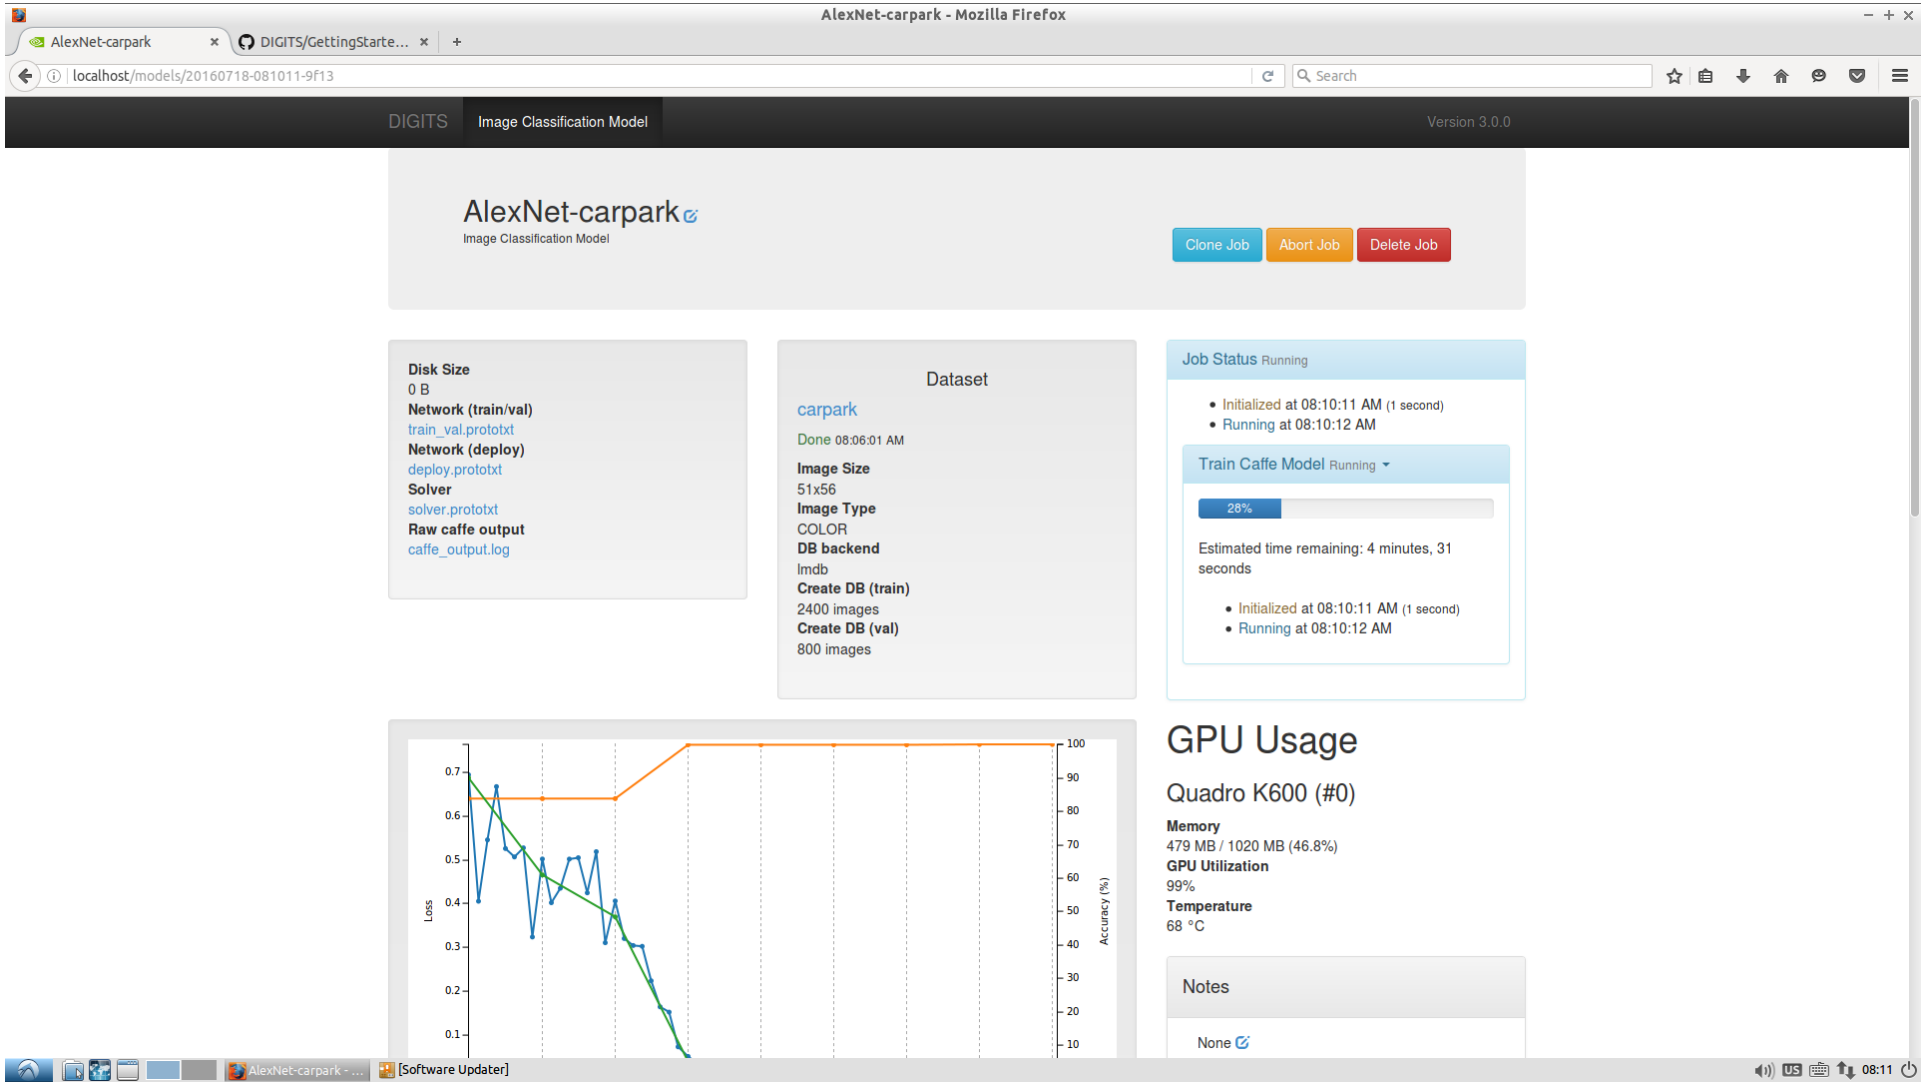

Digits – Testing and Validation

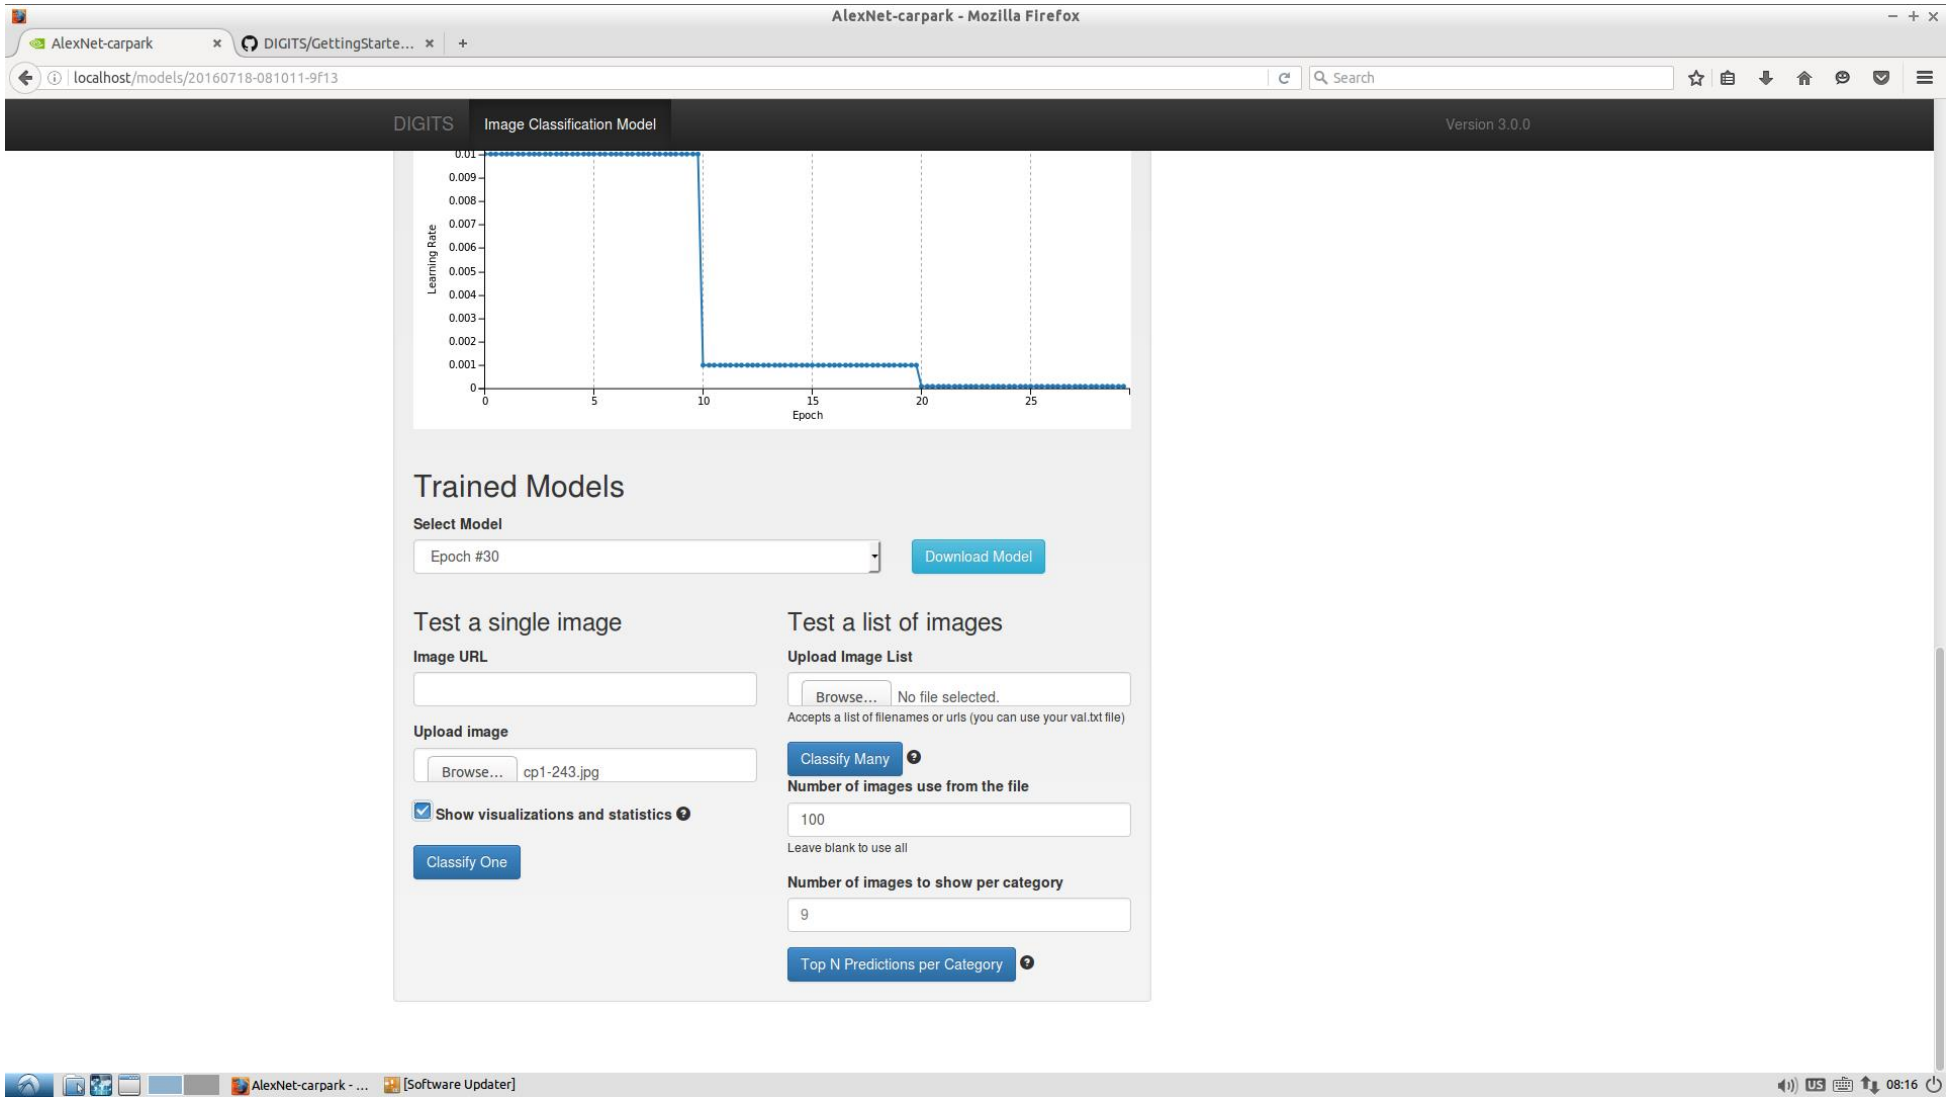

Digits – Results

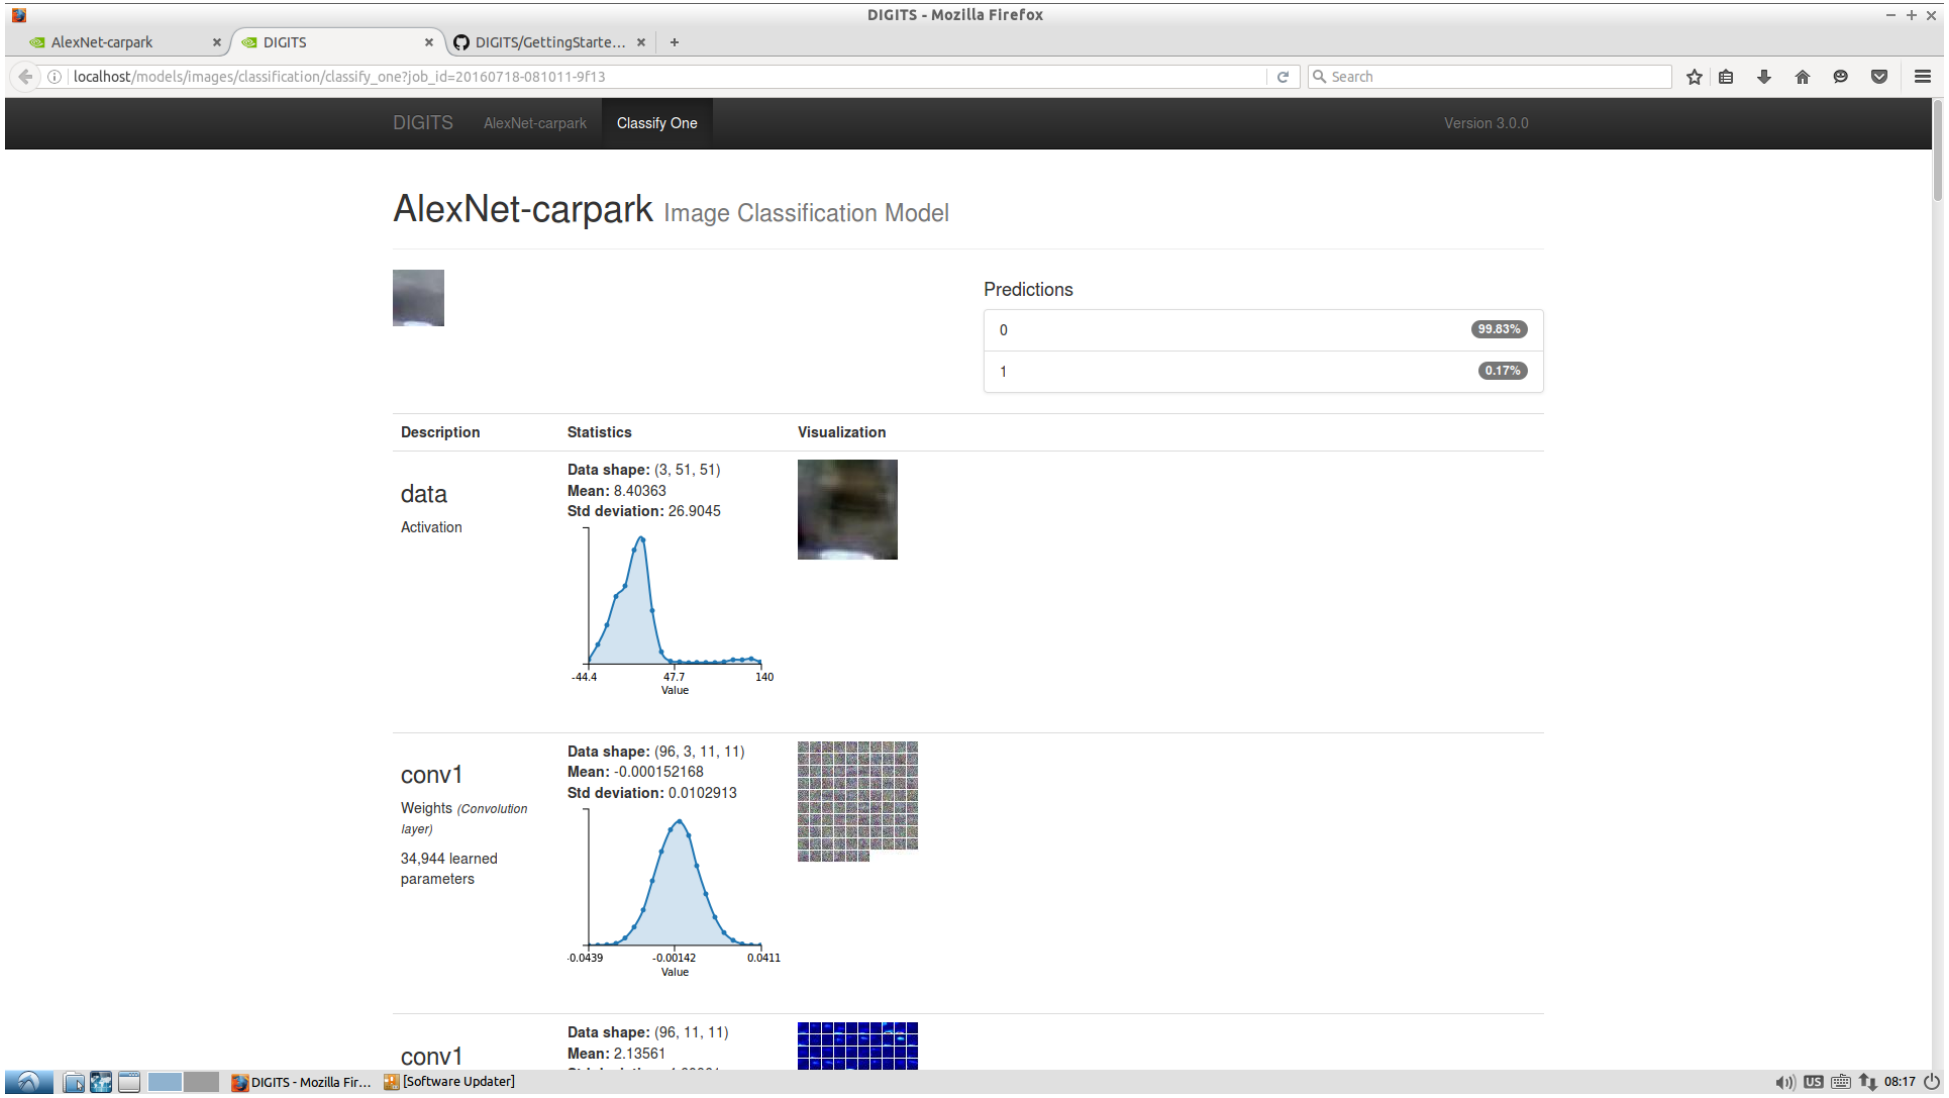

Digits – Results

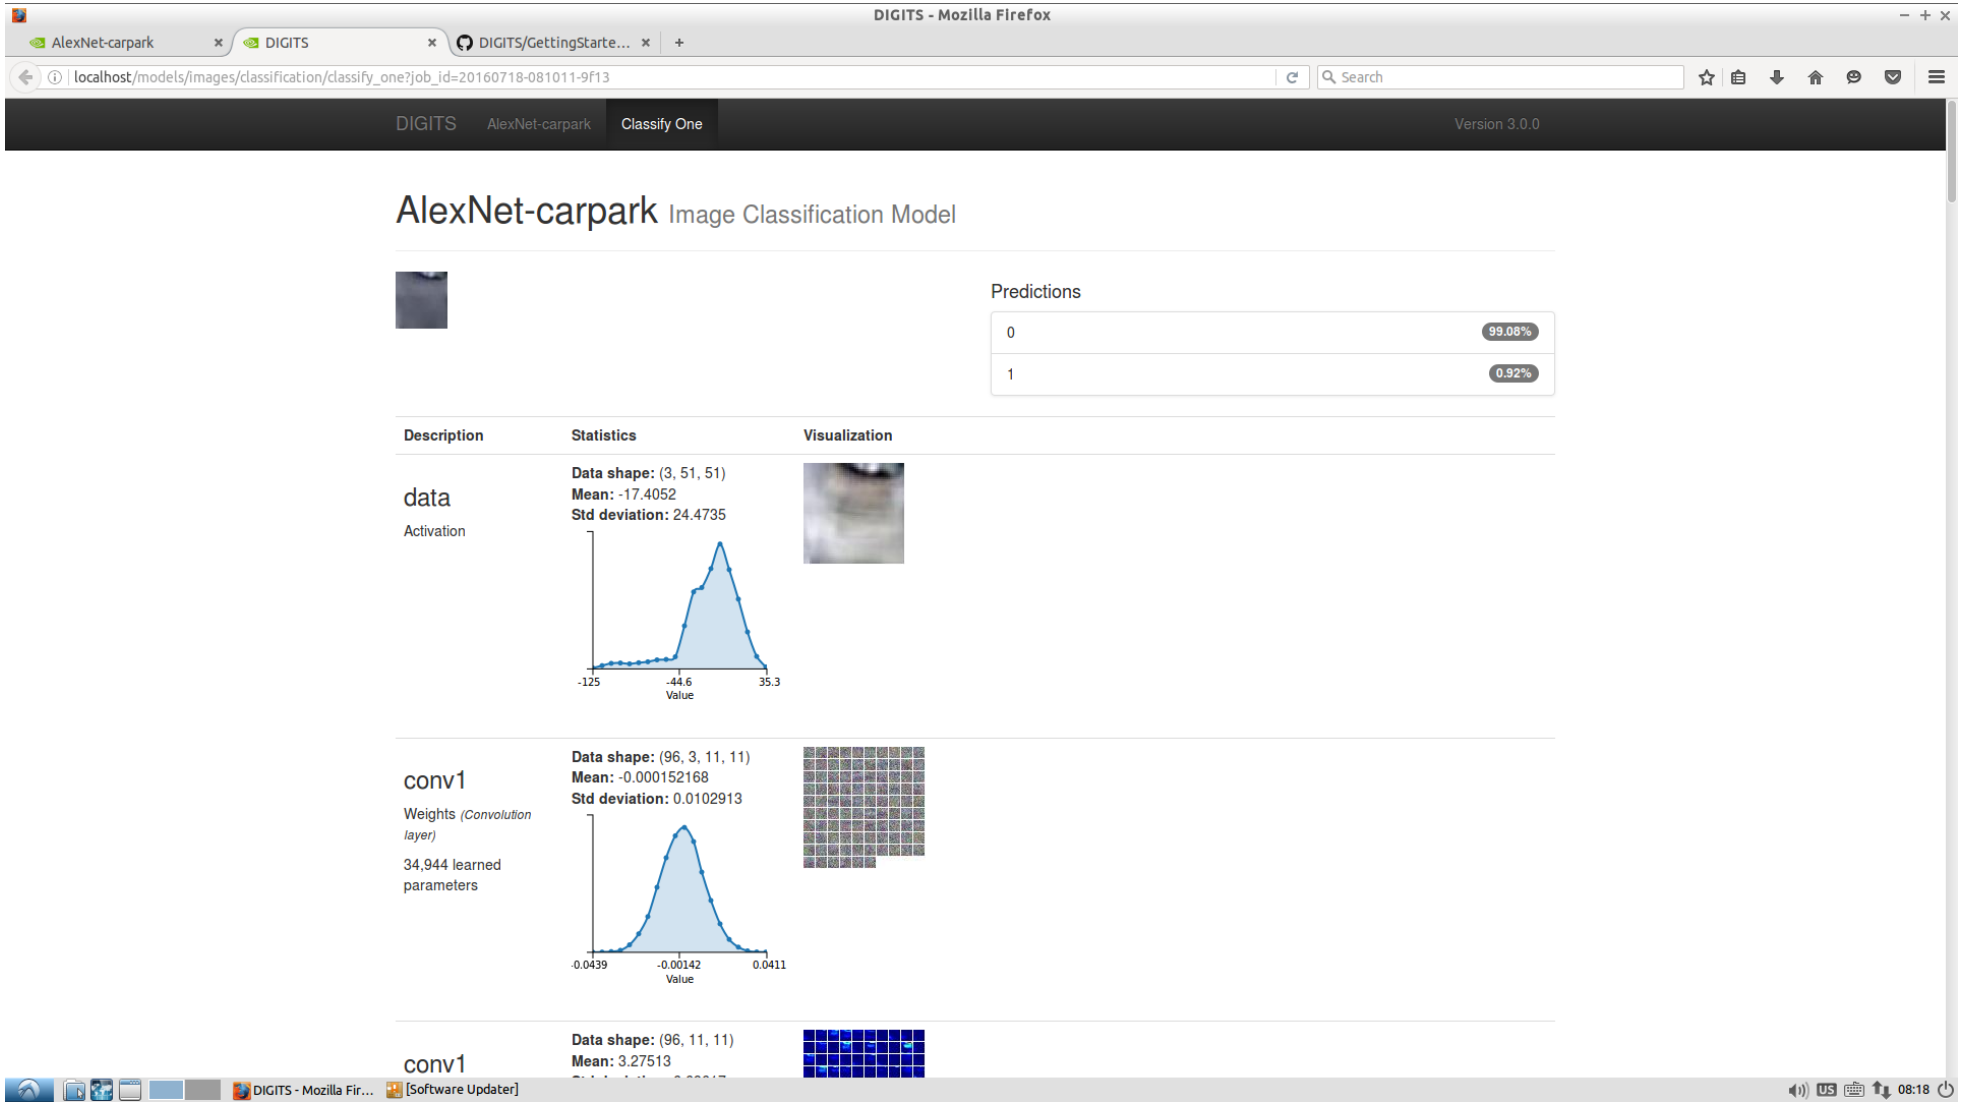

Digits – Results

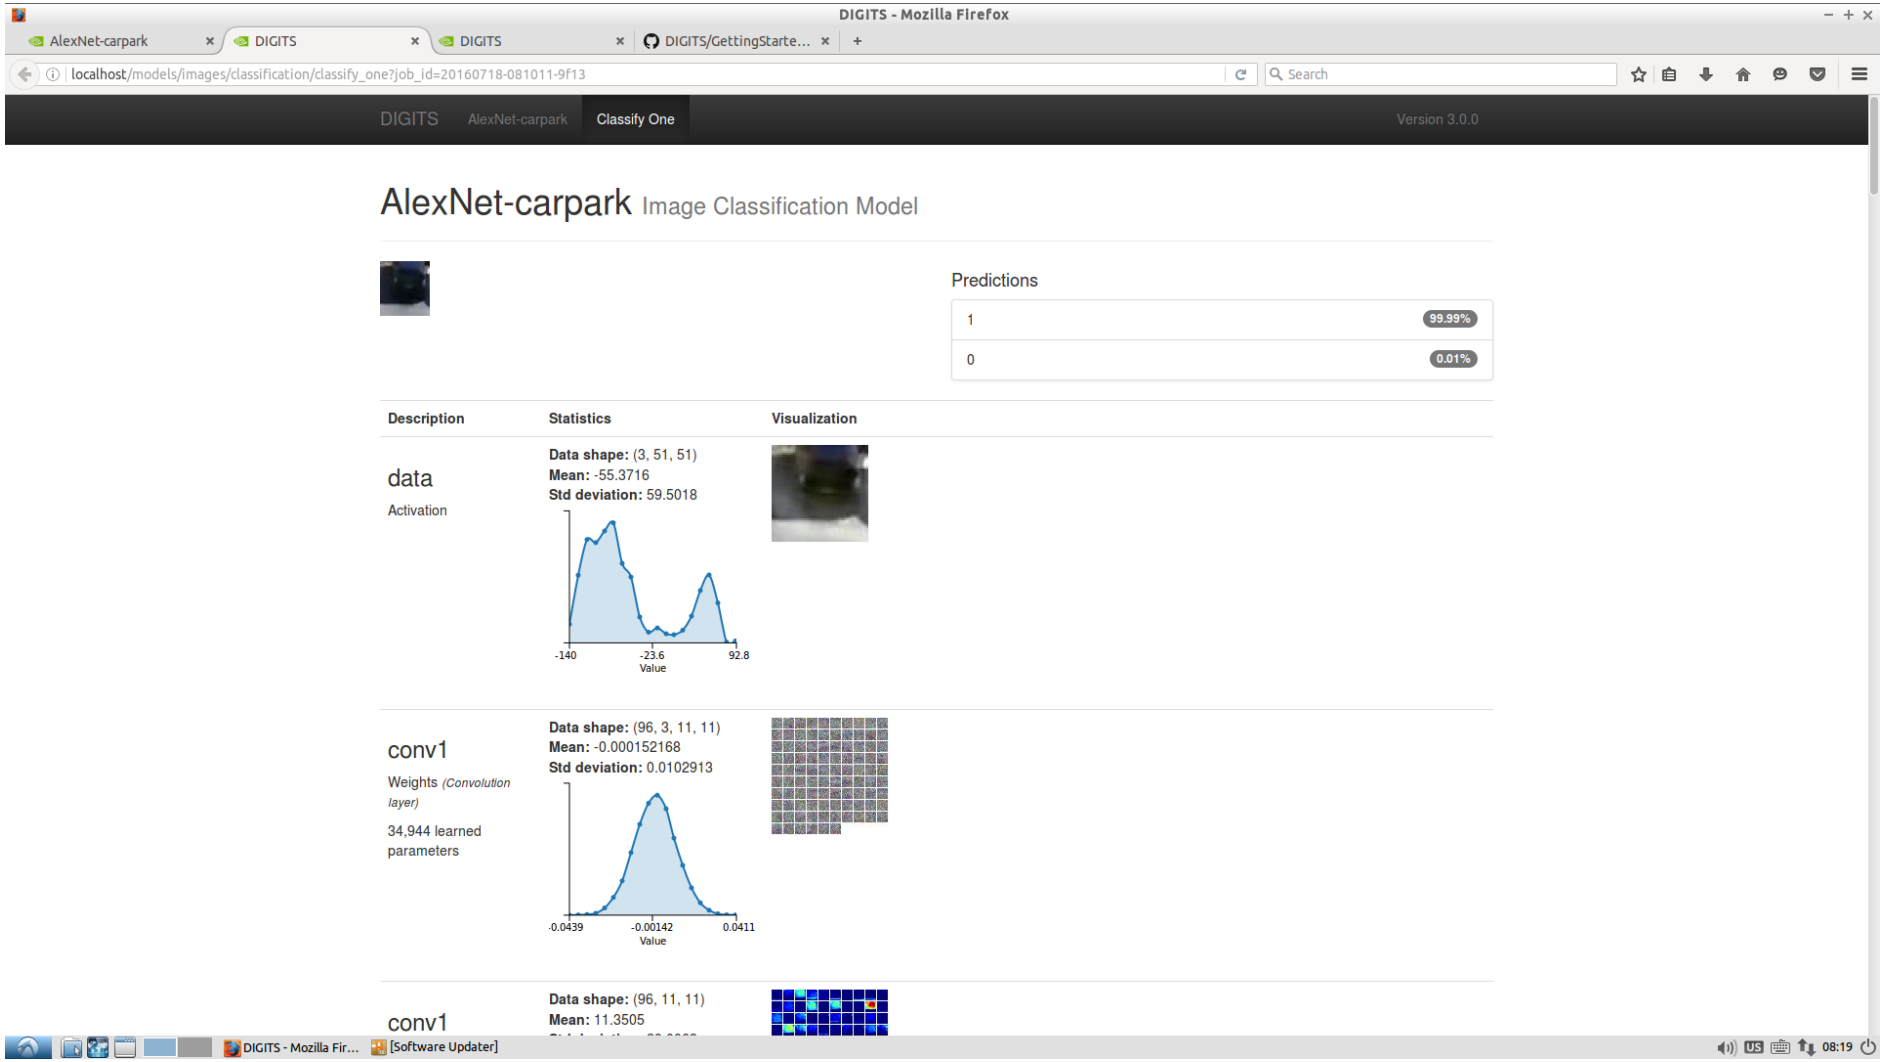

Digits – Further Evaluation

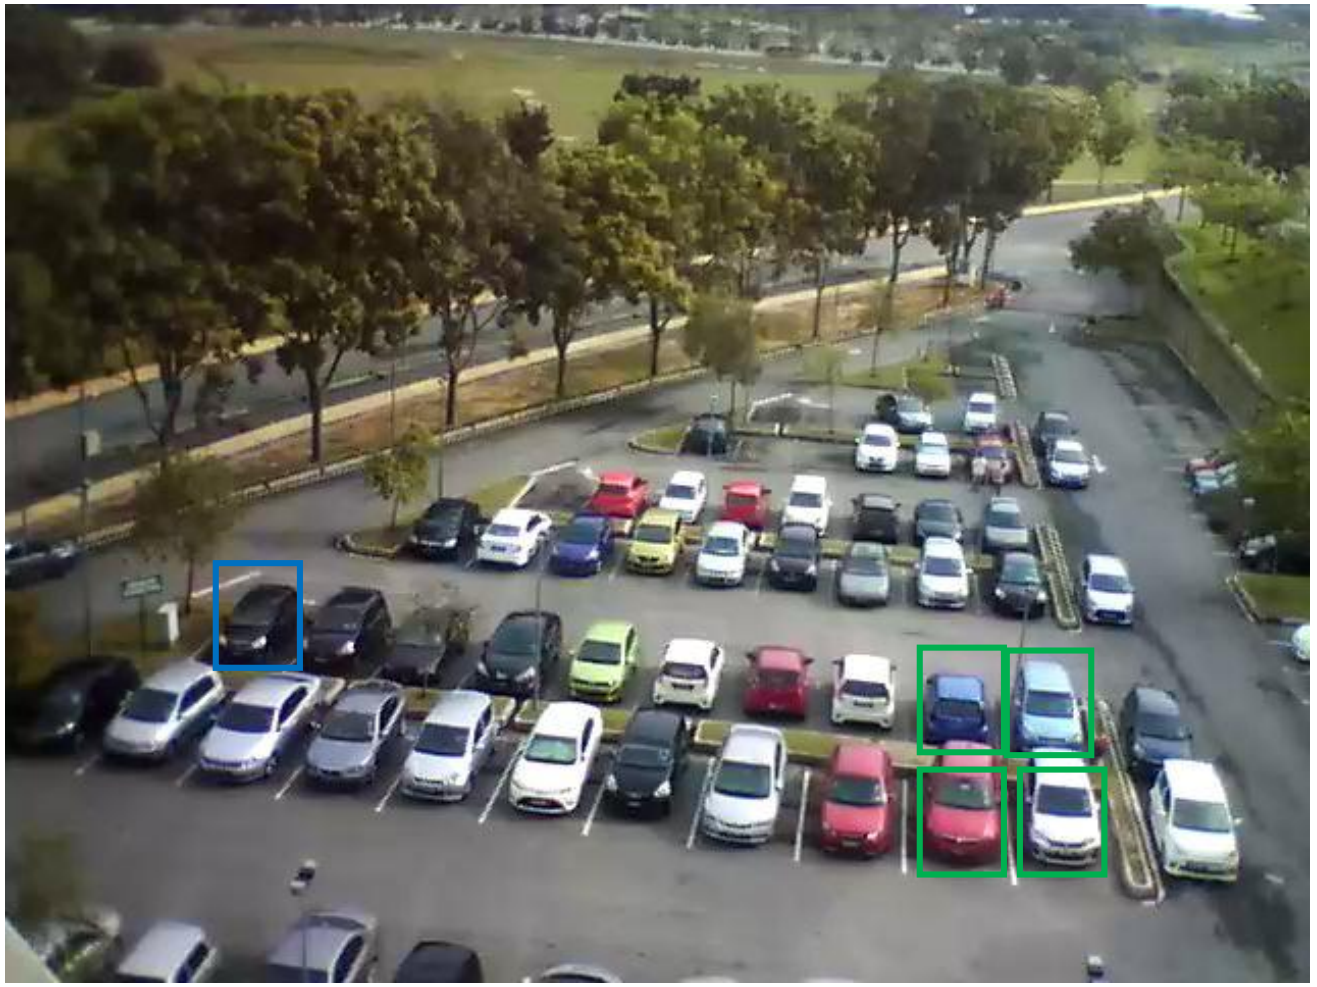

Digits – Further Evaluation

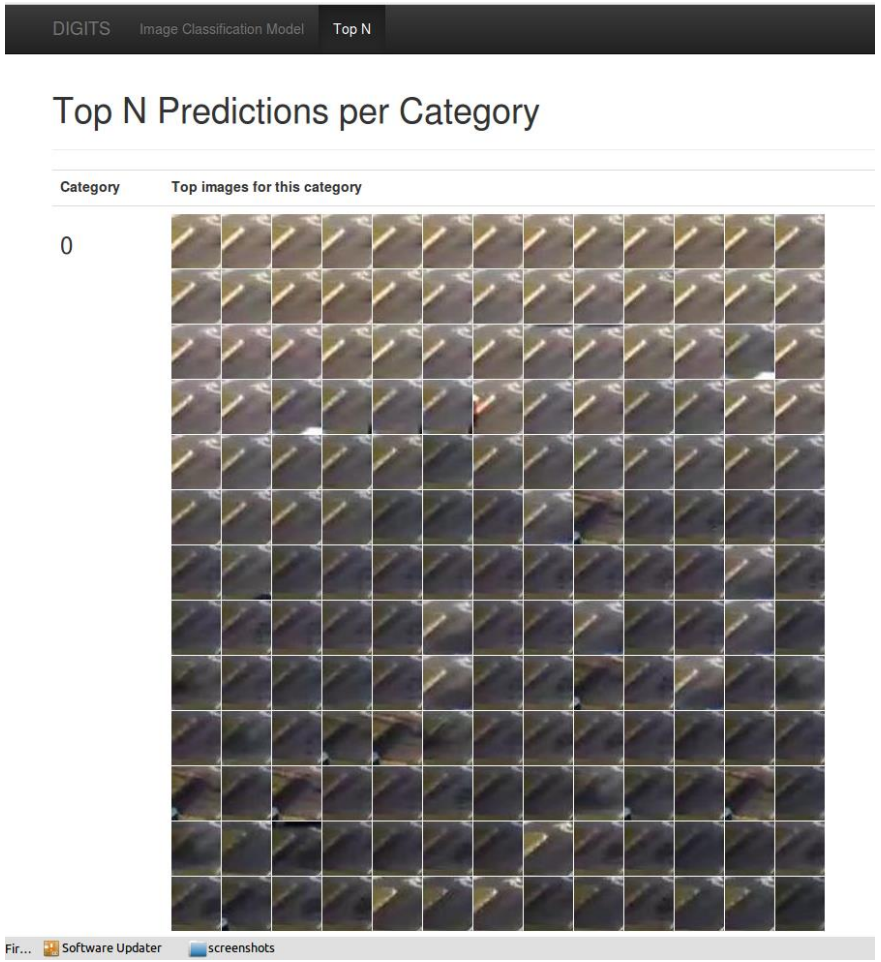

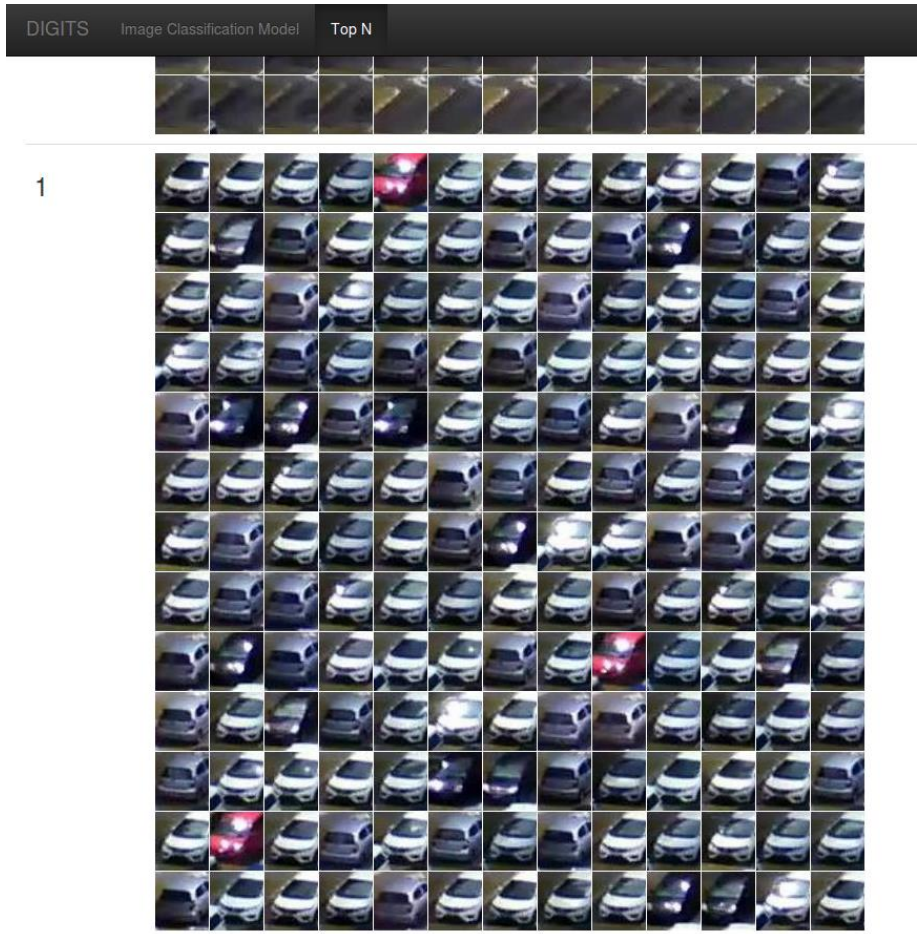

Digits – Further Evaluation

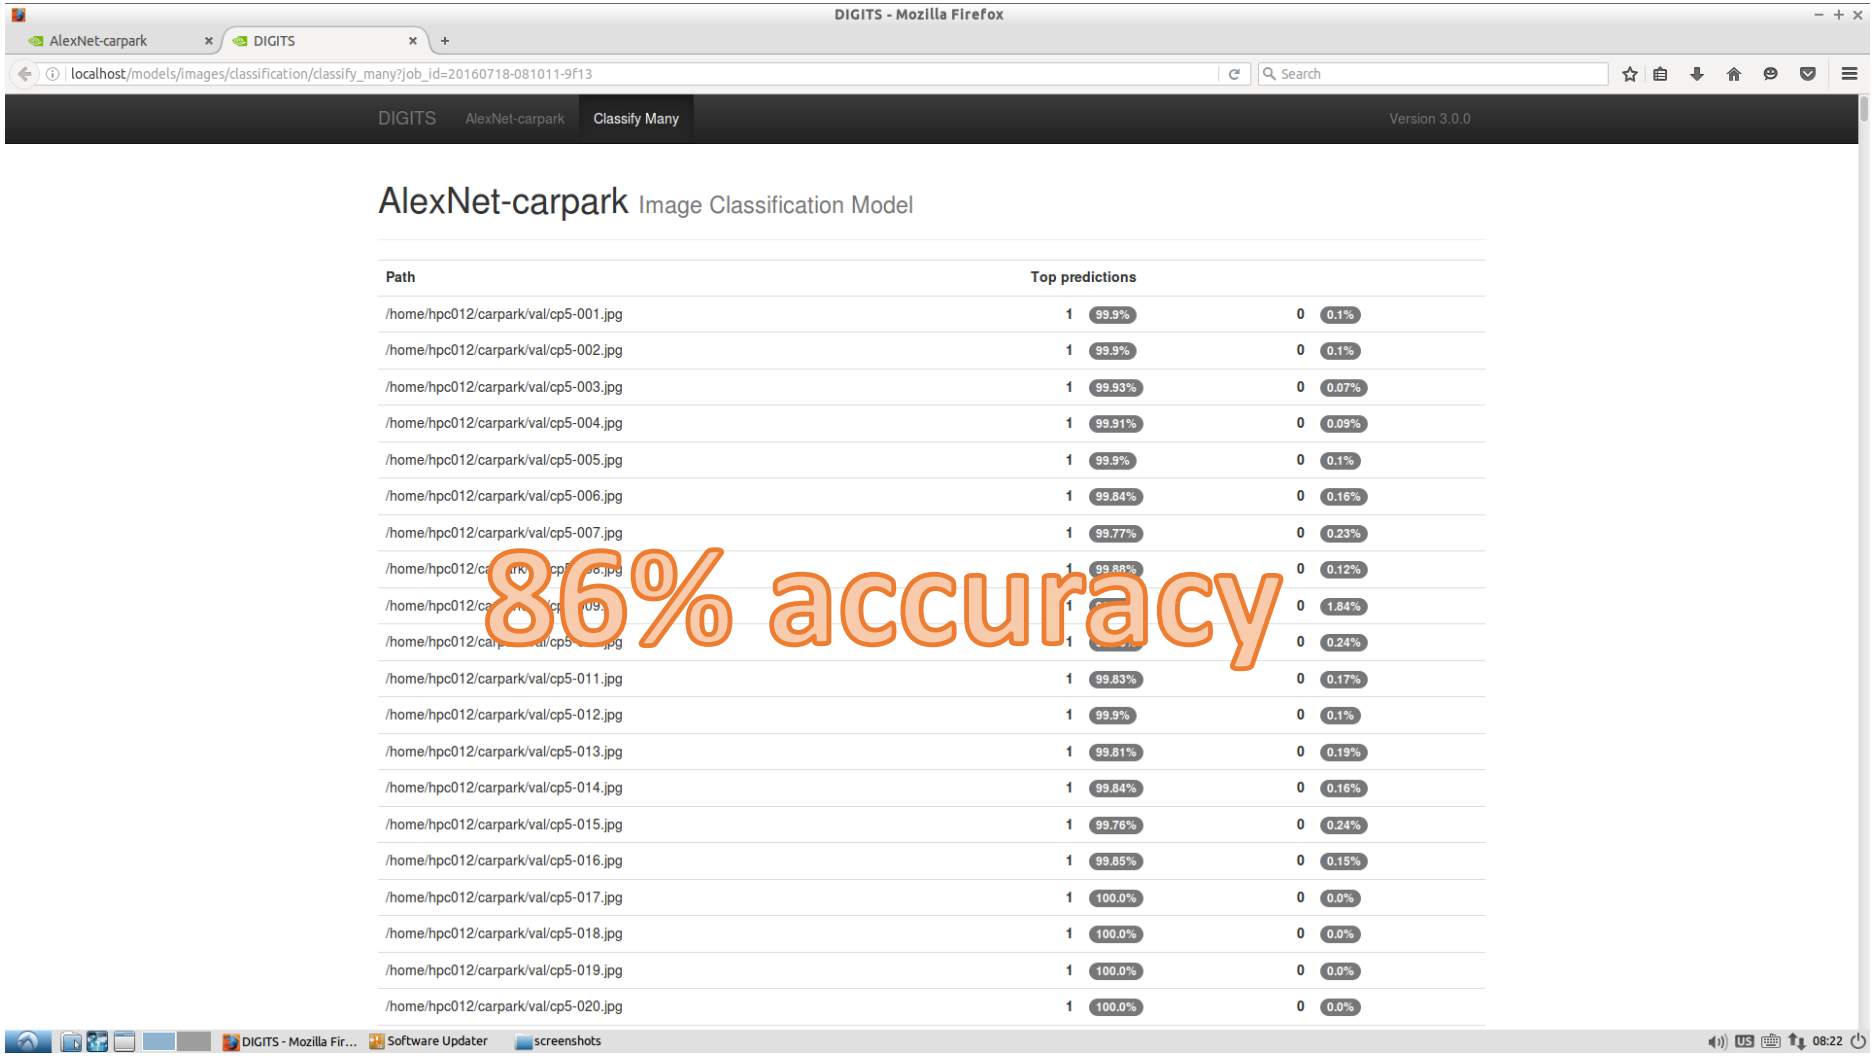

Digits – try it out yourself

- https://github.com/NVIDIA/DIGITS/blob/master/docs/GettingStarted.md

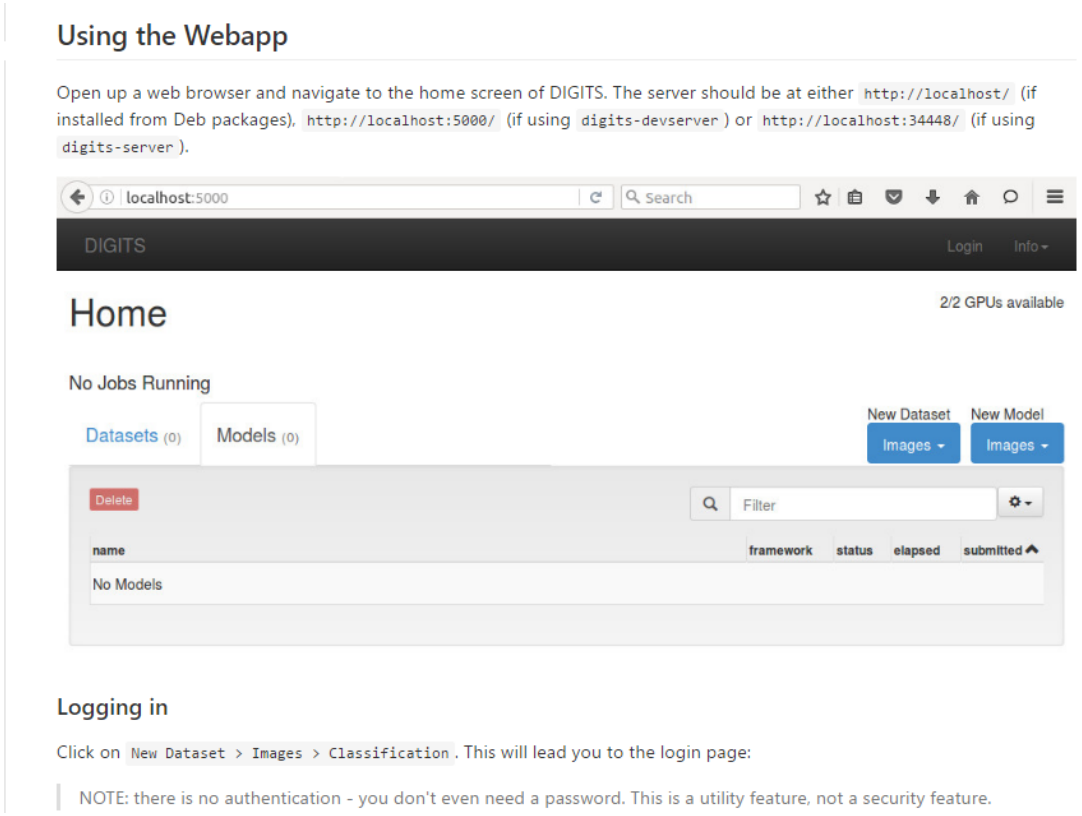

Using the Webapp

Open up a web browser and navigate to the home screen of DIGITS. The server should be at either `http://localhost/` (if installed from Deb packages), `http://localhost:5000/` (if using `digits-devserver`) or `http://localhost:34448/` (if using `digits-server`).

Logging in

Click on `New Dataset > Images > Classification`. This will lead you to the login page:

> NOTE: there is no authentication - you don't even need a password. This is a utility feature, not a security feature.

Thanks!

Questions?

@kuanhoong

https://www.linkedin.com/in/kuanhoong

khpoo@mmu.edu.my